



# RaBiT: RESIDUAL-AWARE BINARIZATION TRAINING FOR ACCURATE AND EFFICIENT LLMs

Anonymous authors

Paper under double-blind review

## ABSTRACT

Efficient deployment of large language models (LLMs) requires extreme quantization, forcing a critical trade-off between low-bit efficiency and performance. Residual binarization promises hardware-friendly, matmul-free inference by stacking binary ( $\pm 1$ ) layers, but is plagued by pathological feature **co-adaptation**. We identify a key failure mode, which we term **inter-path adaptation**: during Quantization-Aware Training (QAT), parallel residual binary paths learn redundant features, degrading the error-compensation structure and crippling the model’s expressive capacity. While prior work relies on heuristic workarounds (*e.g.*, path freezing) that limit model capacity, we propose **RaBiT**, a novel quantization framework that resolves co-adaptation by algorithmically enforcing a residual hierarchy. Its core mechanism sequentially derives each binary path from a single shared full-precision weight, ensuring each path corrects its predecessor’s error. This process is stabilized by a robust initialization that prioritizes functional preservation over mere weight approximation. RaBiT redefines the 2-bit accuracy-efficiency frontier: it achieves state-of-the-art performance, rivals even hardware-intensive Vector Quantization (VQ) methods, and delivers a **4.49 $\times$  inference speed-up** over full-precision models.

## 1 INTRODUCTION

The massive scale of large language models (LLMs) makes model compression essential for their efficient deployment. While 4-bit quantization methods (Frantar et al., 2023; Lin et al., 2024) have emerged as a successful industry standard (Kwon et al., 2023; Zheng et al., 2024), the relentless pursuit of greater efficiency is pushing the research frontier toward the extreme 2-bit regime. This push toward lower bit compression, however, introduces a critical architectural trade-off that defines the current landscape.

At this frontier, two dominant strategies present a stark choice between accuracy and hardware efficiency. On one hand, Vector Quantization (VQ) methods achieve high accuracy but often introduce hardware overhead from lookup tables or complex rotations (Tseng et al., 2024a;b; Egiazarian et al., 2024). On the other hand, residual binarization—stacking multiple binary layers—offers exceptional, matmul-free efficiency. Yet, this highly efficient approach has consistently struggled to maintain performance, hampered by fundamental training challenges that have prevented it from realizing its full potential (Bulat et al., 2024; Wang et al., 2024; Tran & Nguyen, 2025).

The core promise of a residual architecture—that subsequent paths compensate for the errors of preceding ones—is fundamentally undermined by feature **co-adaptation** (Hinton et al., 2012), a pathological training dynamic where parallel components learn redundant features. In residual binarization, we identify a critical manifestation of this phenomenon, which we term **inter-path adaptation**. During standard Quantization-Aware Training (QAT) (Bengio et al., 2013; Hubara et al., 2018), the structurally agnostic global gradient is applied to all paths simultaneously. This forces them to learn redundant features in a race to minimize the global objective, overriding their intended compensatory roles. The result is a breakdown of the residual hierarchy that severely limits the model’s expressive power.

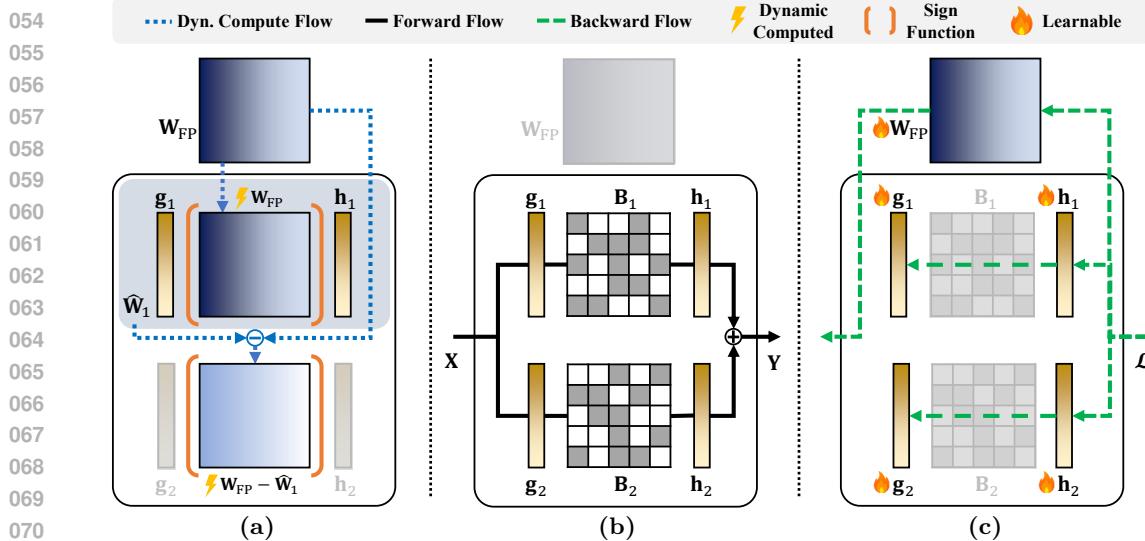


Figure 1: **Overview of the RaBiT training framework.** (a) **Dynamic Compute Process:** During training, binary paths are dynamically derived from a shared weight  $\mathbf{W}_{FP}$  to enforce a residual hierarchy. (b) **Forward Pass:** For inference, these paths execute in parallel for matmul-free efficiency. (c) **Backward Pass:** Gradients from the loss  $\mathcal{L}$  update both the learnable scales ( $\mathbf{g}_i, \mathbf{h}_i$ ) and the shared  $\mathbf{W}_{FP}$ .

Prior attempts to mitigate co-adaptation have relied on heuristic workarounds, such as freezing paths (Bulat et al., 2024; Tran & Nguyen, 2025), which limit the model’s capacity to find an optimal joint solution. To address this, we propose **Residual-Aware Binarization Training (RaBiT)**, a QAT framework that resolves inter-path adaptation by design, as depicted in Figure 1. Instead of using independent latent weights, RaBiT maintains a single shared full-precision weight from which binary paths are sequentially derived on-the-fly, guided by learnable scales. This algorithmically enforces the residual hierarchy, training each path to correct its predecessor’s error. Combined with a robust, function-aware initialization, RaBiT achieves state-of-the-art accuracy while delivering a  $4.49\times$  inference speed-up and halving the training memory footprint.

Our contributions can be summarized as follows:

- We identify and analyze **inter-path adaptation**, a critical manifestation of feature co-adaptation in residual binarization, where the intended error-compensation structure breaks down during Standard QAT as parallel paths become functionally redundant.
- We propose **RaBiT**, a novel QAT framework that resolves inter-path adaptation by enforcing **residual coupling** on-the-fly. The mechanism inherently **halves the training memory footprint** and is stabilized by a robust **function-aware initialization** strategy to tame the unstable dynamics of extreme QAT.
- We demonstrate that RaBiT achieves **state-of-the-art accuracy** at 2-bit precision, delivering a  $4.49\times$  inference speed-up while maintaining competitive performance against hardware-intensive VQ methods through matmul-free operations.

## 2 RELATED WORKS

**The Shift to QAT in Extreme Quantization.** Post-Training Quantization (PTQ) methods, such as GPTQ (Frantar et al., 2023) and AWQ (Lin et al., 2024), have proven highly successful for compressing large language models to 3- or 4-bit precision by focusing on weight approximation. However, these techniques face a steep performance cliff at lower bit-widths (*e.g.*, 2-bit) (Wang et al., 2023), as the information loss from coarse quantization becomes too severe to overcome by simply minimizing weight reconstruction error. Consequently, the research community is shifting from approximating weights to preserving the model’s overall **functionality** (Liu et al., 2025) through Quantization-Aware Training (QAT) (Hubara et al.,

108 2018; Krishnamoorthi, 2018). QAT integrates the simulation of low-precision arithmetic  
 109 into the fine-tuning process, allowing the model to adapt its parameters to the constraints  
 110 of the target bit-width. While QAT is challenging due to the non-differentiable nature of  
 111 quantization—typically addressed with the Straight-Through Estimator (STE) (Bengio et al.,  
 112 2013)—modern frameworks for binary models have found stability by maintaining a latent  
 113 full-precision weight for training and updating it via a surrogate gradient (Wang et al., 2023;  
 114 Xu et al., 2024; Jo et al., 2024; Lee et al., 2025). Our work builds on this robust method to  
 115 address the unique challenges of residual binary architectures.

116 **Co-adaptation in Residual Binary Architectures.** To enhance the limited expressive  
 117 capacity of a single low-bit layer, residual binarization stacks multiple low-bit paths  
 118 ( $\mathbf{W} \approx \sum_i \hat{\mathbf{W}}_i$ ) to achieve higher precision while retaining matmul-free efficiency (Wang  
 119 et al., 2024). However, this parallel architecture is highly susceptible to **feature co-**  
 120 **adaptation** (Hinton et al., 2012), a training pathology where components learn redundant  
 121 features. This phenomenon, which spurred the development of regularization techniques like  
 122 Dropout (Srivastava et al., 2014).

123 In the context of residual binarization, we identify a critical form of feature co-adaptation,  
 124 termed **inter-path adaptation**, where a shared QAT gradient forces parallel paths to  
 125 learn redundant features, undermining their error-compensation hierarchy. While prior work  
 126 relied on suboptimal heuristics like path freezing (Bulat et al., 2024; Tran & Nguyen, 2025)  
 127 that preclude finding a joint optimal solution, RaBiT resolves this core challenge by design,  
 128 enabling true joint optimization while algorithmically enforcing the hierarchy.

### 130 3 MOTIVATION

131 The goal of Quantization-Aware Training (QAT) is to make a quantized student model,  $\mathbf{Y}_s$ ,  
 132 functionally mimic its full-precision teacher,  $\mathbf{Y}_t$ . This is typically achieved by optimizing an  
 133 objective that combines the final task loss with an intermediate knowledge distillation loss,  
 134 often formulated as the mean squared error (MSE) (Hinton et al., 2015; Liu et al., 2024). While  
 135 our full training objective also includes the final KL divergence-based task loss, we focus  
 136 our analysis on the MSE component for its analytical tractability. The additive structure  
 137 of the MSE provides a clear window into how parallel paths interact. In a 2-bit residual  
 138 architecture, the MSE between the teacher output  $y_t$  and the student output  $y_s = y_1 + y_2$   
 139 can be decomposed. Using the Pearson correlation coefficient,<sup>1</sup>, this decomposition is:

$$141 \text{MSE}(y_t, y_s) = \underbrace{(\mathbb{E}[y_t^2] + \mathbb{E}[y_1^2] + \mathbb{E}[y_2^2] - 2\mathbb{E}[y_t y_s])}_{C'} + \underbrace{2\sigma_1\sigma_2}_{\text{PathAmp.}} \cdot \underbrace{\text{Corr}(y_1, y_2)}_{\text{PathCorr.}}$$

142 where  $C'$  represents the sum of correlation-independent error terms. This reveals a core  
 143 principle: to minimize the MSE, the paths must be strongly **negatively correlated**. A  
 144 negative correlation transforms the interaction term into a substantial **bonus** that actively  
 145 reduces the total loss, signifying effective error-cancellation.

146 However, Standard QAT structurally fails to achieve this. The shared global gradient  
 147 induces **inter-path adaptation**, forcing both paths to learn redundant features instead  
 148 of their intended compensatory roles (see Appendix A.1). To provide a concrete analysis,  
 149 we decompose the MSE loss for representative layers of Llama2-7B, selected to show the  
 150 characteristics across the early, middle, and late stages of the network, in Table 1.

151 The analysis in Table 1 is definitive. Across early, middle, and late stages of the network,  
 152 the base error term  $C'$  and the path amplitude  $2\sigma_1\sigma_2$  remain comparable between both  
 153 methods. The critical difference lies in the correlation. Standard QAT yields a correlation  
 154 close to zero, resulting in a negligible interaction term that fails to meaningfully reduce the  
 155 total error. In stark contrast, **RaBiT** structurally enforces a strong negative correlation  
 156 (*e.g.*, -0.50 in layer 5) (see Appendix A.1). This transforms the interaction term into a  
 157 significant loss-reducing bonus, systematically lowering the total MSE (see Appendix A.6).

158 <sup>1</sup>The relationship  $\mathbb{E}[y_1 y_2] \approx \sigma_1 \sigma_2 \text{Corr}(y_1, y_2)$  relies on a zero-mean assumption for the path  
 159 outputs. We empirically verify this, finding the omitted  $\mathbb{E}[y_1] \mathbb{E}[y_2]$  term is less than 1% of the  
 160 covariance term and thus negligible.

162  
**Table 1: Detailed Decomposition of MSE Loss across Representative Layers of**  
**Llama2-7B.** The table breaks down the total MSE into its core components. While the  
163 base error ( $C'$ ) and path amplitude ( $2\sigma_1\sigma_2$ ) are comparable, RaBiT consistently generates a  
164 strong negative correlation, creating a significant loss-reducing **Bonus**. In contrast, Standard  
165 QAT’s weak correlation provides a negligible benefit, demonstrating RaBiT’s structural  
166 advantage in error correction.  
167

Layer	Method	Base Error ( $C'$ )	Path Amp. ( $2\sigma_1\sigma_2$ )	Path Corr. (Corr)	Covariance (Amp. $\times$ Corr)	Total MSE ( $C' + \text{Cov.}$ )
Layer 5 (Early)	Standard QAT	0.0019	0.0030	-0.0752	-0.0002	0.0017
	<b>RaBiT (Ours)</b>	0.0023	0.0028	<b>-0.4961</b>	<b>-0.0014</b>	<b>0.0009</b>
Layer 15 (Mid)	Standard QAT	0.0182	0.0214	-0.1240	-0.0026	0.0156
	<b>RaBiT (Ours)</b>	0.0163	0.0200	<b>-0.3418</b>	<b>-0.0068</b>	<b>0.0094</b>
Layer 25 (Late)	Standard QAT	0.0575	0.0728	-0.1279	-0.0093	0.0482
	<b>RaBiT (Ours)</b>	0.0609	0.0801	<b>-0.3535</b>	<b>-0.0283</b>	<b>0.0327</b>

175  
176 This principled enforcement of anti-correlation creates a more stable optimization landscape,  
177 leading to better generalization and superior performance.  
178

## 179 4 METHOD

180  
181 We introduce **RaBiT**, a novel QAT framework that prevents interference between the parallel  
182 paths of stacked binary architectures. To achieve this, RaBiT enforces a clear error-correction  
183 role for each path using a novel **coupled training** loop, and stabilizes the process with a  
184 **function-aware initialization** strategy. An overview of the RaBiT training framework is  
185 illustrated in Figure 1.  
186

### 187 4.1 THE RESIDUAL BINARIZATION ARCHITECTURE

188  
189 To achieve low-bit precision (*e.g.*, 2-bit) while maximizing computational efficiency, we adopt  
190 a residual architecture built upon highly efficient binary building blocks.  
191

**Binary Building Blocks.** The fundamental component is the dual-scale binarization  
192 framework. We define the approximation of a weight matrix  $\hat{\mathbf{W}}$  using a notation that  
193 highlights the underlying element-wise scaling operations:  
194

$$\hat{\mathbf{W}} = \mathbf{g} \odot \mathbf{B} \odot \mathbf{h}. \quad (1)$$

195 Here, the expression denotes an operation where each element of the resulting matrix,  $(\hat{\mathbf{W}})_{ij}$ ,  
196 is computed as  $g_i B_{ij} h_j$ .  $\mathbf{B} \in \{-1, +1\}^{d_{\text{out}} \times d_{\text{in}}}$  is the binary core matrix, and  $\mathbf{g} \in \mathbb{R}^{d_{\text{out}}}$ ,  
197  $\mathbf{h} \in \mathbb{R}^{d_{\text{in}}}$  are full-precision, per-channel scaling vectors. The key advantage of this formulation  
198 is its matmul-free efficiency. For an input vector  $\mathbf{x} \in \mathbb{R}^{d_{\text{in}}}$ , the forward operation computes  
199 the output vector  $\mathbf{y} \in \mathbb{R}^{d_{\text{out}}}$  as  $\mathbf{y} = \mathbf{g} \odot (\mathbf{B}(\mathbf{h} \odot \mathbf{x}))$ , which can be implemented using only  
200 additions and subtractions, eliminating costly multiplications.  
201

202  
203 **Multi-bit Approximation via Stacking.** To enhance representational capacity (*e.g.*, to  
204 2-bit) while retaining this efficiency, we stack  $k = 2$  binary paths in parallel. The effective  
205 weight is the sum of two binarized terms:  
206

$$\hat{\mathbf{W}}^{(k)} = \sum_{i=1}^k \hat{\mathbf{W}}_i = \sum_{i=1}^k \mathbf{g}_i \odot \mathbf{B}_i \odot \mathbf{h}_i. \quad (2)$$

207 This architecture preserves the underlying matmul-free execution, as the forward pass simply  
208 accumulates the outputs from each path.  
209

### 210 4.2 COUPLED TRAINING FOR CO-ADAPTATION MITIGATION

211  
212 To address inter-path adaptation, RaBiT abandons the standard approach of training  
213 independent latent weights for each binary path. Instead, it maintains a **single shared**  
214 **full-precision (FP) weight**  $\mathbf{W}_{\text{FP}}$  that serves as the anchor for the entire residual structure.  
215

**The Coupled Forward Pass.** The core of RaBiT lies in its dynamic forward pass. For a 2-bit architecture with  $k = 2$  paths, the binary core matrices,  $\mathbf{B}_1$  and  $\mathbf{B}_2$ , are not stored but re-calculated during every forward pass from the shared weight  $\mathbf{W}_{\text{FP}}$  (Figure 1a). This process algorithmically enforces the error-compensation hierarchy. Unlike the dynamically derived binary cores, the scaling vectors  $\{\mathbf{g}_i, \mathbf{h}_i\}$  are independent, learnable parameters that capture the magnitude of each path. The derivation proceeds as follows:

1. **Path 1 Derivation:** The first binary core,  $\mathbf{B}_1$ , is determined by directly binarizing the shared weight:  $\mathbf{B}_1 = \text{sign}(\mathbf{W}_{\text{FP}})$ . This binary core is then combined with its corresponding learnable scaling vectors,  $\mathbf{g}_1$  and  $\mathbf{h}_1$ , to reconstruct the first-path approximation,  $\hat{\mathbf{W}}_1 = \mathbf{g}_1 \odot \mathbf{B}_1 \odot \mathbf{h}_1$ .
2. **Residual Calculation:** The residual error,  $\mathbf{R}_1$ , is calculated by subtracting the *reconstructed* first path from the shared weight:  $\mathbf{R}_1 = \mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_1$ .
3. **Path 2 Derivation:** The second binary core,  $\mathbf{B}_2$ , is then determined by binarizing this freshly computed residual error:  $\mathbf{B}_2 = \text{sign}(\mathbf{R}_1)$ . The final effective weight used in the forward pass is the sum of the two reconstructed paths:  $\hat{\mathbf{W}}^{(2)} = \hat{\mathbf{W}}_1 + (\mathbf{g}_2 \odot \mathbf{B}_2 \odot \mathbf{h}_2)$ .

A key design choice is to derive only the **binary cores**  $\mathbf{B}_i := \text{sign}(\mathbf{R}_{i-1})$  dynamically, while treating the **scaling vectors**  $\{\mathbf{g}_i, \mathbf{h}_i\}$  as independent, learnable parameters. This separation of roles is crucial for both computational efficiency and training stability. Re-calculating optimal scales for the residual at every forward pass—*e.g.*, via Singular Value Decomposition (SVD)—would be prohibitively expensive. Consequently, by maintaining them as learnable parameters, the optimizer can leverage state accumulation (*e.g.*, momentum) to robustly fine-tune the well-initialized values (Section 4.3). This data-adaptive tuning is vital for training stability and allows the error-compensation hierarchy to function effectively by learning the optimal magnitude for each path.

**Backward Pass and Parameter Updates.** The backward pass is designed for stability and effectiveness. The gradient from the loss  $\mathcal{L}$  flows back to update both the independent, learnable scaling vectors  $\{\mathbf{g}_i, \mathbf{h}_i\}$  and the single shared weight  $\mathbf{W}_{\text{FP}}$ , as shown in Figure 1c.

- **Gradient for Learnable Scales:** The scaling vectors  $\{\mathbf{g}_i, \mathbf{h}_i\}$  are treated as standard learnable parameters and receive their own gradients via the chain rule. For a mini-batch of size  $B$ , the gradients are accumulated over each sample, treating the dynamic binary cores ( $\mathbf{B}_i$ ) as constants during this calculation:

$$\nabla_{\mathbf{g}_i} = \sum_{b=1}^B \Delta_b \odot (\mathbf{B}_i (\mathbf{h}_i \odot \mathbf{X}_b)), \quad \nabla_{\mathbf{h}_i} = \sum_{b=1}^B (\mathbf{B}_i^\top (\Delta_b \odot \mathbf{g}_i)) \odot \mathbf{X}_b. \quad (3)$$

Here,  $b$  is the sample index within the mini-batch,  $\mathbf{X}_b$  is the input vector for that sample, and  $\Delta_b = (\partial \mathcal{L} / \partial \mathbf{Y}_b)$  is the upstream gradient from the layer’s output  $\mathbf{Y}_b$  for that sample.

- **Gradient for the Shared Weight:** To update the single shared weight  $\mathbf{W}_{\text{FP}}$ , RaBiT uses an **effective-weight gradient**. This acts as a Straight-Through Estimator (STE) for the *entire coupled derivation process*. The gradient is computed with respect to the final effective weight  $\hat{\mathbf{W}}^{(k)} = \sum_i \hat{\mathbf{W}}_i$  and is passed back directly to update  $\mathbf{W}_{\text{FP}}$ :

$$\nabla_{\mathbf{W}_{\text{FP}}} \approx \nabla_{\hat{\mathbf{W}}^{(k)}} \mathcal{L} = (\partial \mathcal{L} / \partial \mathbf{Y})^\top \mathbf{X}. \quad (4)$$

In this context,  $\mathcal{L}$  is the task loss, while  $\mathbf{X}$  and  $\mathbf{Y}$  represent the full input and output matrices for the mini-batch. This update completes the training loop: by recomputing the binary paths from the updated  $\mathbf{W}_{\text{FP}}$  at every step, RaBiT continuously forces each path to correct the latest residual error, which preserves the overall error-compensation hierarchy.

For inference, the final binary cores  $\{\mathbf{B}_i\}$  are derived from the trained  $\mathbf{W}_{\text{FP}}$  and then frozen, while the shared weight  $\mathbf{W}_{\text{FP}}$  is discarded. The resulting architecture is highly efficient, as the independent binary paths execute in a fully parallel, matmul-free manner. Crucially, this single-weight design also provides a key training advantage: by halving the latent parameters, it reduces the memory required for optimizer states by 50%, a major bottleneck in LLM fine-tuning. The complete training step is detailed in Algorithm 2.

270 4.3 STABLE INITIALIZATION FOR FUNCTIONAL PRESERVATION  
271

272 QAT in the 2-bit regime is extremely sensitive to the initial quantization error. To mitigate this, we propose a two-stage initialization process that prioritizes preserving model  
273 functionality over merely approximating weight values.  
274

275 **1. Iterative Residual SVID.** The core of our initialization is to find a set of binary  
276 paths that jointly approximate a target weight matrix. A standard greedy decomposition is  
277 suboptimal because the choice for the first path irreversibly biases all subsequent paths. To  
278 find a better joint solution, we propose **Iterative Residual Sign-Value-Independent**  
279 **Decomposition (SVID)**, a Gauss-Seidel style iteration that allows the paths to co-adapt.  
280 The process iteratively refines the scales  $\{\mathbf{g}_i, \mathbf{h}_i\}$  and binary cores  $\{\mathbf{B}_i\}$  for each path  
281  $i = 1, \dots, k$  over  $t = 1, \dots, T$  iterations as follows:  
282

$$\begin{cases} \mathbf{R}_i^{(t)} &:= \mathbf{W}_{\text{FP}} - \left( \sum_{j < i} \hat{\mathbf{W}}_j^{(t)} + \sum_{j > i} \hat{\mathbf{W}}_j^{(t-1)} \right), \\ \mathbf{B}_i^{(t)}, \mathbf{g}_i^{(t)}, \mathbf{h}_i^{(t)} &:= \text{SVID}(\mathbf{R}_i^{(t)}), \\ \hat{\mathbf{W}}_i^{(t)} &:= \mathbf{g}_i^{(t)} \odot \mathbf{B}_i^{(t)} \odot \mathbf{h}_i^{(t)}. \end{cases} \quad (5)$$

283 Here,  $\text{SVID}(\cdot)$  (Xu et al., 2024) finds optimal per-channel scales by separating the signs and  
284 performing a rank-1 SVD approximation on the magnitudes. Note that this iterative process  
285 is not directly applied to the raw weights  $\mathbf{W}_{\text{FP}}$ , but to a preconditioned target matrix  $\mathbf{W}'$ ,  
286 which we define next.  
287

288 **2. I/O Channel Importance-Scaled Preconditioning.** To ensure our iterative de-  
289 composition focuses on the most functionally critical components of the weight matrix, we  
290 do not apply it to the raw weights  $\mathbf{W}_{\text{FP}}$ . Instead, inspired by recent work on preserving  
291 functional saliency (Boža & Hradiš, 2025), we first **precondition** the matrix to create the  
292 target  $\mathbf{W}'$ . Using a small calibration dataset, we compute input activation magnitudes ( $\mathbf{s}_{\text{in}}$ )  
293 and output gradient magnitudes ( $\mathbf{s}_{\text{out}}$ ) and re-weight the full-precision matrix accordingly:  
294

$$\mathbf{W}' = \mathbf{s}_{\text{out}}^{\alpha_{\text{out}}} \odot \mathbf{W}_{\text{FP}} \odot \mathbf{s}_{\text{in}}^{\alpha_{\text{in}}}. \quad (6)$$

295 Finally, after running the iterative SVID, the resulting scales are mapped back to the original  
296 domain for training:  $\mathbf{g}_i = \mathbf{s}_{\text{out}}^{-\alpha_{\text{out}}} \odot \mathbf{g}'_i$  and  $\mathbf{h}_i = \mathbf{s}_{\text{in}}^{-\alpha_{\text{in}}} \odot \mathbf{h}'_i$ . This strategy dramatically  
297 reduces the initial task loss, ensuring a stable start to QAT (see Algorithm 1 for the full  
298 algorithm, with analysis in Table 6 and Figure 5).  
299

300 5 EXPERIMENTS  
301302 5.1 EXPERIMENTAL SETTINGS  
303

304 **Setup.** We evaluate RaBiT on recent LLMs including Llama2/3 and Gemma3 (Touvron  
305 et al., 2023; AI@Meta, 2024). For QAT, we use a 200M-token subset from a combined  
306 WikiText-2 and C4 dataset (Jo et al., 2024). We report perplexity (PPL) on their validation  
307 sets (context length: 4096) and the average zero-shot accuracy (QA Avg.) on five com-  
308 mon sense reasoning benchmarks (*e.g.*, HellaSwag, PIQA) (Sakaguchi et al., 2021; Zellers  
309 et al., 2019; Clark et al., 2018; Bisk et al., 2020), with a detailed breakdown provided in  
310 Appendix A.3.  
311

312 **Training Details.** We employ a QAT framework with knowledge distillation (KD) (Hinton  
313 et al., 2015; Liu et al., 2024), where the full-precision model serves as the teacher. The  
314 objective function combines Kullback–Leibler (KL) divergence loss on the output logits with  
315 intermediate MSE losses:  $\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{kl}} + \gamma \sum_i \mathcal{L}_{\text{inter},i}$ , with  $\gamma = 10$  for Llama-family, but  
316  $\gamma = 0$  for Gemma3 models to avoid instability from their large activation range. All models  
317 are trained for 6 epochs with the Muon optimizer (Jordan et al., 2024) and our proposed  
318 function-aware initialization. Full hyperparameters are listed in Appendix A.5.  
319

320 **Baselines.** We benchmark RaBiT against a comprehensive set of state-of-the-art 2- to 3-bit  
321 methods. Baselines include (1) standard methods like GPTQ and EfficientQAT (Frantar et al.,  
322

Table 2: **Comparison with state-of-the-art 2-3-bit methods on Llama models.** We report perplexity (PPL  $\downarrow$ ) and zero-shot QA Average ( $\uparrow$ ). For the 2-bit results, the best and runner-up are marked in **bold** and underlined, respectively. RaBiT achieves state-of-the-art (SOTA) performance on Llama2-7B and Llama3-8B, while showing highly competitive results on Llama2-13B.

Methods	Llama-2-7B					Llama-2-13B					Llama-3-8B					
	bit	Wiki2 $\downarrow$	C4 $\downarrow$	QA Avg $\uparrow$	bit	Wiki2 $\downarrow$	C4 $\downarrow$	QA Avg $\uparrow$	bit	Wiki2 $\downarrow$	C4 $\downarrow$	QA Avg $\uparrow$	bit	Wiki2 $\downarrow$	C4 $\downarrow$	
Baseline	16	5.12	6.63	62.26	16	4.57	6.05	65.46	16	5.75	8.32	68.66				
GPTQ	2.1	50.75	36.76	39.16	2.1	43.84	23.07	43.72	2	1.21e3	4.97e2	35.59				
EfficientQAT	2.1	6.42	8.34	57.75	2.1	5.58	7.40	62.07	2.1	8.75	12.09	60.63				
AQLM	2.3	6.29	8.56	58.57	2.2	5.41	7.20	61.58	2.3	7.23	10.32	64.12				
QuIP#	2	6.19	8.16	58.23	2	5.35	7.20	61.96	2	8.70	12.04	63.89				
QTIP	2	<u>5.86</u>	<u>7.73</u>	<u>58.97</u>	2	<b>5.11</b>	<b>6.85</b>	<b>62.92</b>	2	<u>7.52</u>	<u>10.76</u>	<u>63.88</u>				
BitStack	3	6.91	9.10	56.54	3	5.90	7.86	61.06	3	12.38	17.51	58.41				
	2	29.97	34.91	40.12	2	67.98	72.60	39.38	2	2.75e3	1.93e3	36.21				
DB-LLM	2	7.23	9.62	55.12	2	6.19	8.38	59.41	2	12.08	16.80	50.92				
MBOK	3	6.13	8.13	54.63	3	5.14	6.94	62.73	3	7.81	11.29	61.08				
	2	6.99	9.38	53.63	2	5.76	7.89	60.58	2	10.74	14.61	54.41				
DBF	2.3	5.81	7.69	59.84	2.3	5.15	6.85	62.53	2.3	7.22	10.34	64.84				
	2	6.10	8.05	58.42	2	5.33	7.13	61.53	2	7.78	10.99	62.90				
RaBiT(Ours)	3	5.36	7.06	63.05	3	4.84	6.51	64.09	3	6.58	9.54	65.61				
	2	<b>5.77</b>	<b>7.64</b>	<b>61.51</b>	2	<u>5.15</u>	<u>6.95</u>	<u>62.10</u>	2	<b>7.34</b>	<b>10.52</b>	<b>64.13</b>				

Table 3: **Comparison with state-of-the-art 2-bit methods on Gemma3 models.** We report perplexity (PPL  $\downarrow$ ) and zero-shot QA Average ( $\uparrow$ ). The context length is 4096. RaBiT consistently achieves SOTA or highly competitive performance, demonstrating its robustness across diverse model architectures.

Methods	Gemma3-1B					Gemma3-4B					Gemma3-12B					
	bit	Wiki2 $\downarrow$	C4 $\downarrow$	QA Avg $\uparrow$	bit	Wiki2 $\downarrow$	C4 $\downarrow$	QA Avg $\uparrow$	bit	Wiki2 $\downarrow$	C4 $\downarrow$	QA Avg $\uparrow$	bit	Wiki2 $\downarrow$	C4 $\downarrow$	
Baseline	16	9.80	13.69	57.82	16	6.88	10.44	67.60	16	5.50	9.28	73.45				
DBF	2	13.28	17.57	51.98	2	8.72	12.71	60.91	2	6.97	10.60	68.37				
QTIP	2	13.14	17.36	50.30	2	8.31	12.21	<b>63.47</b>	2	<b>6.65</b>	10.25	<b>69.69</b>				
RaBiT(Ours)	2	<b>11.27</b>	<b>15.54</b>	<b>53.18</b>	2	<b>8.09</b>	<b>11.91</b>	62.21	2	6.66	<b>10.18</b>	68.85				

2023; Huang et al., 2024); (2) high-accuracy but hardware-intensive Vector Quantization (VQ) approaches such as AQLM, QuIP#, and QTIP (Egiazarian et al., 2024; Tseng et al., 2024a;b); and (3) hardware-efficient binary/residual methods like BitStack, DB-LLM, MBOK, and DBF (Wang et al., 2024; Chen et al., 2024; Tran & Nguyen, 2025; Boža & Hradiš, 2025)<sup>2</sup>, which are the most direct architectural competitors.

## 5.2 MAIN RESULTS

As shown in Tables 2 and 3, RaBiT consistently redefines the state-of-the-art for 2-bit quantization, demonstrating superior performance across all tested models.

**Dominance over Hardware-Efficient Methods.** RaBiT significantly outperforms other matmul-free binary/residual methods. On Llama2-7B, its 5.77 WikiText-2 PPL is a marked improvement over competitors like MBOK (6.99 PPL) and DBF (6.10 PPL). This performance gap widens on larger models and more complex datasets, underscoring the severe performance penalty incurred by the **inter-path adaptation** that these methods fail to address. The catastrophic failure of some methods, such as BitStack on Llama3-8B (2.75e3 PPL), highlights the instability that RaBiT’s principled design successfully overcomes.

**Achieving VQ-Level Accuracy with Binary Efficiency.** More impressively, RaBiT closes the gap with and often surpasses hardware-intensive VQ methods, resolving the historical trade-off between accuracy and efficiency. On Llama2-7B, RaBiT’s 5.77 PPL edges out the leading VQ method, QTIP (5.86 PPL), setting a new SOTA for 2-bit quantization.

<sup>2</sup>For DB-LLM and MBOK, we used our re-implementation

This trend holds for reasoning tasks, where RaBiT’s 61.51% QA average on Llama2-7B surpasses all other listed methods, including QTIP’s 58.97%, demonstrating superior functional preservation. RaBiT’s robustness is further evident on Llama3-8B, where it maintains strong performance (7.34 PPL) while VQ methods like QuIP# suffer from severe degradation (8.70 PPL), showcasing the stability of our training framework.

### 384 5.3 ABLATION STUDIES

#### 385 5.3.1 COMPONENT-WISE CONTRIBUTION ANALYSIS

388 Channel Importance Scaling (**S**) also yields  
 389 a consistent, albeit smaller, improvement by  
 390 prioritizing salient weights. The full RaBiT  
 391 model, integrating all components, achieves  
 392 the optimal 5.77 PPL, demonstrating that a  
 393 residually-coupled training mechanism com-  
 394 bined with function-aware initialization is es-  
 395 sential for state-of-the-art performance. We  
 396 performed an ablation study to analyze the  
 397 contributions of RaBiT’s core components:  
 398 Coupled QAT, Iterative SVID (**I**), and I/O  
 399 Channel Importance-Scaled Preconditioning  
 (**S**), with results in Table 4.

400 The analysis clearly shows that **Coupled QAT** is the most critical performance factor.  
 401 Simply switching from Standard QAT (6.55 PPL) to Coupled QAT drops the perplexity to  
 402 5.83, confirming that resolving inter-path adaptation yields the largest gain.

403 Our initialization methods (**I** and **S**) provide further essential improvements. While they  
 404 offer a significant boost to the baseline Standard QAT, their role within the powerful Coupled  
 405 QAT framework is to provide the final, crucial fine-tuning needed to reach the optimal 5.77  
 406 PPL. This synergy between a robust training method and a function-aware initialization is  
 407 key to RaBiT’s state-of-the-art performance.

#### 409 5.3.2 ANALYSIS OF COUPLED TRAINING DYNAMICS

411 To empirically validate that coupled training resolves co-adaptation, we conducted a controlled  
 412 experiment comparing RaBiT to four variants: (1) **Standard QAT** (independent latent  
 413 weights), (2) **MBOK** (frozen primary binary core), (3) **Scale-only** (frozen binary cores),  
 414 and (4) **Scale-frozen** (RaBiT with frozen scales). All variants shared the same initialization  
 415 and hyperparameters; for a fair comparison of training dynamics, MBOK also used the same  
 416 optimizer as our model.

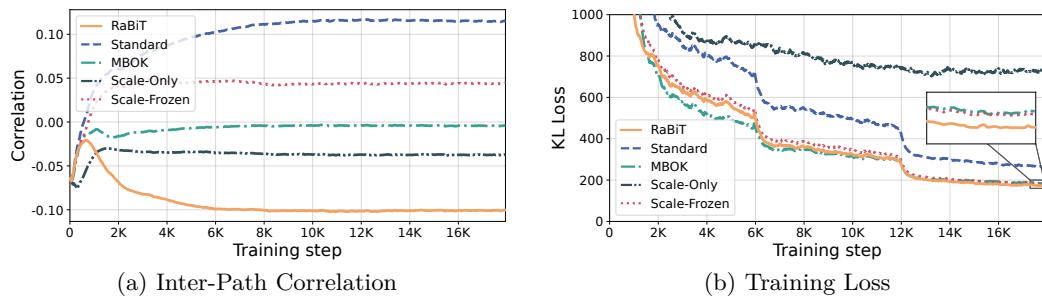
417 Figure 2 reveals the resulting training dynamics. As theorized, **RaBiT** successfully maintains  
 418 a stable negative inter-path correlation, enforcing the error-correction hierarchy (Figure 2a).  
 419 In contrast, **Standard QAT** develops a strong positive correlation, confirming that a shared  
 420 global gradient induces harmful redundancy. The constrained variants (**MBOK**, **Scale-**  
 421 **frozen**) fail to establish a strong anti-correlation, limiting their optimization potential. This  
 422 structural advantage directly translates to model functionality, as shown by the training loss  
 423 curves (Figure 2b). RaBiT achieves the lowest and most stable loss, while the co-adaptation  
 424 in Standard QAT and the incomplete optimization of the other variants lead to significantly  
 425 higher loss. This analysis confirms that RaBiT’s ability to jointly optimize all parameters  
 426 while algorithmically preventing co-adaptation is the key to its superior performance.

### 427 5.4 INFERENCE PERFORMANCE

429 RaBiT not only achieves SOTA accuracy but also delivers exceptional inference speed by  
 430 leveraging its parallelizable matmul-free binary architecture. As shown in Table 5, RaBiT  
 431 achieves up to a **4.49 × speed-up** in end-to-end decoding throughput for a 256-token  
 generation over the FP16 baseline on an NVIDIA RTX 4090.

387 **Table 4: Ablation on RaBiT** (Llama2-7B  
 388 PPL). The analysis isolates the impact of  
 389 Iterative Residual SVID (**I**) and I/O Channel  
 390 Importance-Scaled Preconditioning (**S**).

Training Method	I	S	Wiki2 ↓
Standard QAT	✓		6.55
		✓	6.21
	✓	✓	6.31
<b>Coupled QAT</b> (RaBiT)	✓		5.83
		✓	5.78
	✓	✓	5.81
			<b>5.77</b>



**Figure 2: Visualization of Coupled Training Dynamics.** (a) **Inter-Path Correlation:** RaBiT enforces a negative inter-path correlation, indicating effective error-correction, whereas Standard QAT leads to positive correlation (co-adaptation). (b) **Training Loss:** This structural advantage directly translates to a lower and more stable training loss for RaBiT, demonstrating its superior optimization path.

This performance gain stems from two key advantages. First, the  $8 \times$  reduction in model size (2-bit vs. 16-bit) dramatically lowers memory bandwidth requirements, which is the primary bottleneck in the autoregressive decoding phase. Second, unlike VQ methods, RaBiT avoids hardware-unfriendly overheads like lookup tables or rotations. Its simple architecture of additions and element-wise scaling translates to higher hardware utilization. This is evident in our kernel-level benchmarks, where RaBiT’s specialized kernels exhibit consistently lower latency than both the FP16 baseline and QTIP’s VQ kernels. By eliminating computational complexity, RaBiT ensures that theoretical memory savings translate directly into real-world speed, delivering a solution that is both accurate and genuinely efficient. Further details on our kernel design and additional performance benchmarks are provided in Appendix A.4.1 and Appendix A.4.2, respectively.

**Table 5: Inference Performance Analysis on NVIDIA RTX 4090.** Kernel latency for key Llama2-7B/13B layers and Llama2-7B decoding throughput for a 256-token generation. **RaBiT shows superior efficiency** at both the kernel and system levels.

Method	bit	Kernel-Level Latency ( $\mu\text{s}$ ) ↓				End-to-End Decoding Throughput (tok/s) ↑
		4096×4096 (q_proj, 7B)	11008×4096 (gate_proj, 7B)	5120×5120 (q_proj, 13B)	13824×5120 (gate_proj, 13B)	
FP16	16	17.15 (1.00×)	70.37 (1.00×)	17.85 (1.00×)	122.90 (1.00×)	64.96 (1.00×)
DBF	2.3	12.66 (1.35×)	28.43 (2.48×)	14.72 (1.21×)	31.87 (3.86×)	157.66 (2.43×)
	2	11.47 (1.50×)	20.90 (3.37×)	14.08 (1.27×)	29.58 (4.15×)	175.21 (2.70×)
QTIP	3	24.04 (0.71×)	37.08 (1.90×)	36.22 (0.49×)	49.97 (2.46×)	153.59 (2.36×)
	2	23.40 (0.73×)	42.40 (1.66×)	37.46 (0.48×)	59.22 (2.08×)	171.74 (2.64×)
RaBiT (Ours)	3	<b>8.15 (2.10×)</b>	<b>17.13 (4.11×)</b>	<b>9.90 (1.80×)</b>	<b>22.36 (5.50×)</b>	<b>191.63 (2.95×)</b>
	2	<b>7.72 (2.22×)</b>	<b>15.71 (4.48×)</b>	<b>8.33 (2.14×)</b>	<b>17.50 (7.02×)</b>	<b>291.88 (4.49×)</b>

## 6 CONCLUSION

This paper resolves the critical trade-off between accuracy and hardware efficiency in 2-bit LLM quantization by introducing **RaBiT**. We first identify and analyze **inter-path adaptation** as a fundamental bottleneck that compromises the error-compensation structure in residual binarization. RaBiT’s core mechanism, on-the-fly **residual coupling**, algorithmically prevents this breakdown during training, ensuring the model’s full expressive power is utilized. To tame the notoriously unstable dynamics of extreme quantization, we further introduce a robust **function-aware initialization** strategy that ensures stable convergence. Our comprehensive experiments demonstrate that RaBiT achieves new **state-of-the-art performance** at 2-bit precision, surpassing not only existing binary methods but also complex and hardware-intensive vector quantization approaches. By establishing this new frontier, RaBiT paves the way for the efficient deployment of high-performance low-bit LLMs and provides a scalable foundation for future research.

---

**486 REPRODUCIBILITY STATEMENT**

488 We are committed to ensuring the reproducibility of our research. To this end, we provide de-  
 489 tailed descriptions of our methodology, experimental setup, and hyperparameters throughout  
 490 the paper and its appendices.

492 **Algorithm and Implementation Details.** While full source code availability is subject  
 493 to our organization’s review protocols, we have made every effort to ensure reproducibility by  
 494 providing detailed algorithmic descriptions. The core logic of the **RaBiT** QAT framework is  
 495 thoroughly explained in Section 4 and presented as step-by-step pseudocode in Algorithm 2.  
 496 Similarly, the design and optimization principles of our high-performance CUDA kernel are  
 497 described in Appendix A.4.1, providing a clear blueprint for implementation.

498 **Data.** Our training process utilizes a 200M-token subset of the publicly available WikiText-  
 499 2 and C4 datasets, generated using the same data processing approach as described in Jo  
 500 et al. (2024)<sup>3</sup>. All evaluations are performed on standard public benchmarks (WikiText-2,  
 501 C4, HellaSwag, PIQA, ARC-e, ARC-c, and WinoGrande), as detailed in Section 5.1.  
 502

503 **Hyperparameters and Infrastructure.** A comprehensive list of all hyperparameters  
 504 used for training each model, including model-specific learning rates, optimizer settings, and  
 505 the initialization parameters ( $\alpha_{\text{in}}, \alpha_{\text{out}}, T_{\max}$ ), is provided in Appendix A.5. All experiments  
 506 were conducted on a single node equipped with four NVIDIA H100 GPUs. The reported  
 507 results are from a single training run for each model, which is a standard practice for LLM  
 508 fine-tuning at this scale due to the high computational cost.

**509 REFERENCES**

- 510 AI@Meta. The Llama 3 herd of models. arXiv:2407.21783, 2024.
- 513 Yoshua Bengio, Nicholas Léonard, and Aaron Courville. Estimating or propagating gradients  
 514 through stochastic neurons for conditional computation. arXiv:1308.3432, 2013.
- 515 Yonatan Bisk, Rowan Zellers, Jianfeng Gao, and Yejin Choi. PIQA: Reasoning about physical  
 516 commonsense in natural language. In *Proceedings of the AAAI Conference on Artificial  
 517 Intelligence (AAAI)*, 2020.
- 519 Vladimír Boža and Michal Hradiš. Addition is almost all you need: Compressing neural  
 520 networks with double binary factorization. arXiv:2505.11076, 2025.
- 521 Adrian Bulat, Yassine Ouali, and Georgios Tzimiropoulos. QBB: Quantization with binary  
 522 bases for llms. *Advances in Neural Information Processing Systems*, 37:3209–3228, 2024.
- 524 Hong Chen, Chengtao Lv, Chaofan Yang, Yuxuan Zhang, Zhong-Qiu Wang, Heng Zhang, Jie  
 525 Chen, Enhong Chen, and Dacheng Tao. DB-LLM: Accurate dual-binarization for efficient  
 526 LLMs. In *Findings of the Association for Computational Linguistics (ACL Findings)*,  
 527 2024.
- 528 Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick,  
 529 and Oyvind Tafjord. Think you have solved question answering? Try ARC, the AI2  
 530 Reasoning Challenge. In *Proceedings of the 2018 Conference on Empirical Methods in  
 531 Natural Language Processing (EMNLP)*, 2018.
- 532 Vage Egiazarian, Andrei Panferov, Denis Kuznedelev, Elias Frantar, Artem Babenko, and  
 533 Dan Alistarh. Extreme compression of large language models via additive quantization.  
 534 *arXiv preprint arXiv:2401.06118*, 2024.
- 535 Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. GPTQ: Accurate post-  
 536 training quantization for generative pre-trained transformers. In *Advances in Neural  
 537 Information Processing Systems (NeurIPS)*, 2023.

---

539 <sup>3</sup>The specific code we used is available in the official repository: <https://github.com/dongwonjo/BinaryMoS/blob/main/utils/datautils.py>.

- 540 Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network.  
 541 arXiv:1503.02531, 2015.
- 542
- 543 Geoffrey E Hinton, Nitish Srivastava, Alex Krizhevsky, Ilya Sutskever, and Ruslan R  
 544 Salakhutdinov. Improving neural networks by preventing co-adaptation of feature detectors.  
 545 *arXiv preprint arXiv:1207.0580*, 2012.
- 546
- 547 Hanxian Huang, Zechun Liu, Changsheng Zhao, Yangyang Shi, Raghuraman Krishnamoorthi,  
 548 and Vikas Chandra. EfficientQAT: A block-wise quantization-aware fine-tuning frame-  
 549 work for large language models. In *Advances in Neural Information Processing Systems*  
 (NeurIPS), 2024.
- 550
- 551 Itay Hubara, Matthieu Courbariaux, Daniel Soudry, Ran El-Yaniv, and Yoshua Bengio.  
 552 Quantized neural networks: Training neural networks with low precision weights and  
 553 activations. *journal of machine learning research*, 18(187):1–30, 2018.
- 554
- 555 Dongwon Jo, Taesu Kim, Yulhwa Kim, and Jae-Joon Kim. Mixture of Scales: Memory-  
 556 efficient token-adaptive binarization for large language models. In *Advances in Neural*  
 557 *Information Processing Systems (NeurIPS)*, 2024.
- 558
- 559 Keller Jordan, Yuchen Jin, Vlado Boza, You Jiacheng, Franz Cesista, Laker Newhouse, and  
 560 Jeremy Bernstein. Muon: An optimizer for hidden layers in neural networks, 2024. URL  
<https://kellerjordan.github.io/posts/muon/>.
- 561
- 562 Raghuraman Krishnamoorthi. Quantizing deep convolutional networks for efficient inference:  
 563 A whitepaper. *arXiv preprint arXiv:1806.08342*, 2018.
- 564
- 565 Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu,  
 566 Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large  
 567 language model serving with pagedattention. In *Proceedings of the ACM SIGOPS 29th*  
*Symposium on Operating Systems Principles*, 2023.
- 568
- 569 Banseok Lee, Dongkyu Kim, Youngcheon You, and Youngmin Kim. Littlebit: Ultra low-bit  
 570 quantization via latent factorization. *arXiv preprint arXiv:2506.13771*, 2025.
- 571
- 572 Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangx-  
 573 uan Xiao, Xingyu Dang, Chuang Gan, and Song Han. AWQ: Activation-aware weight  
 574 quantization for on-device LLM compression and acceleration. *Proceedings of Machine*  
*Learning and Systems (MLSys)*, 6:87–100, 2024.
- 575
- 576 Zechun Liu, Barlas Oguz, Changsheng Zhao, Ernie Chang, Pierre Stock, Yashar Mehdad,  
 577 Yangyang Shi, Raghuraman Krishnamoorthi, and Vikas Chandra. LLM-QAT: Data-free  
 578 quantization aware training for large language models. In *Findings of the Association for*  
*Computational Linguistics (ACL Findings)*, 2024.
- 579
- 580 Zechun Liu, Changsheng Zhao, Hanxian Huang, Sijia Chen, Jing Zhang, Jiawei Zhao, Scott  
 581 Roy, Lisa Jin, Yunyang Xiong, Yangyang Shi, et al. ParetoQ: Scaling laws in extremely  
 582 low-bit llm quantization. *arXiv preprint arXiv:2502.02631*, 2025.
- 583
- 584 Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. WinoGrande:  
 585 An adversarial winograd schema challenge at scale. *Communications of the ACM*, 64(9):  
 99–106, 2021.
- 586
- 587 Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdi-  
 588 nov. Dropout: a simple way to prevent neural networks from overfitting. *The journal of*  
*589 machine learning research*, 15(1):1929–1958, 2014.
- 590
- 591 Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei,  
 592 Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2:  
 593 Open foundation and fine-tuned chat models. arXiv:2307.09288, 2023.
- 594
- 595 Ba-Hien Tran and Van Minh Nguyen. Highly efficient and effective llms with multi-boolean  
 596 architectures. *arXiv preprint arXiv:2505.22811*, 2025.

- 594 Albert Tseng, Jerry Chee, Qingyao Sun, Volodymyr Kuleshov, and Christopher De Sa.  
 595 QuIP#: Even better LLM quantization with hadamard incoherence and lattice codebooks.  
 596 In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024a.  
 597
- 598 Albert Tseng, Mark Z Mao, Jerry Chee, Tamas Linder, Volodymyr Kuleshov, Robert M Gray,  
 599 and Christopher De Sa. QTIP: Quantization with trellises and incoherence processing. In  
 600 *Advances in Neural Information Processing Systems (NeurIPS)*, 2024b.
- 601 Hongyu Wang, Shuming Ma, Li Dong, Shaohan Huang, Huaijie Wang, Lingxiao Ma, Fan  
 602 Yang, Ruiping Wang, Yi Wu, and Furu Wei. BitNet: Scaling 1-bit transformers for large  
 603 language models. *arXiv preprint arXiv:2310.11453*, 2023.
- 604 Xinghao Wang, Pengyu Wang, Bo Wang, Dong Zhang, Yunhua Zhou, and Xipeng Qiu.  
 605 Bitstack: Any-size compression of large language models in variable memory environments.  
 606 *arXiv preprint arXiv:2410.23918*, 2024.
- 607
- 608 Yuzhuang Xu, Xu Han, Zonghan Yang, Shuo Wang, Qingfu Zhu, Zhiyuan Liu, Weidong Liu,  
 609 and Wanxiang Che. Onebit: Towards extremely low-bit large language models. *Advances  
 610 in Neural Information Processing Systems*, 37:66357–66382, 2024.
- 611 Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a  
 612 machine really finish your sentence? In *Proceedings of the 57th Annual Meeting of the  
 613 Association for Computational Linguistics (ACL)*, 2019.
- 614
- 615 Feng Zhang, Yanbin Liu, Weihua Li, Jie Lv, Xiaodan Wang, and Quan Bai. Towards superior  
 616 quantization accuracy: A layer-sensitive approach. *arXiv preprint arXiv:2503.06518*, 2025.
- 617 Yang Zhang, Yanfei Dong, and Kenji Kawaguchi. Investigating layer importance in large  
 618 language models. *arXiv preprint arXiv:2409.14381*, 2024.
- 619
- 620 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Livia Sun, Jeff Huang, Cody Hao  
 621 Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E Gonzalez, et al. SGLang:  
 622 Efficient execution of structured language model programs. *Advances in neural information  
 623 processing systems*, 37:62557–62583, 2024.
- 624
- 625
- 626
- 627
- 628
- 629
- 630
- 631
- 632
- 633
- 634
- 635
- 636
- 637
- 638
- 639
- 640
- 641
- 642
- 643
- 644
- 645
- 646
- 647

648    **A APPENDIX**  
 649

650    **A.1 MATHEMATICAL ANALYSIS OF TRAINING DYNAMICS**  
 651

652    This section provides a mathematical analysis of the training dynamics for residual binary  
 653    architectures. We demonstrate why Standard QAT is prone to **inter-path adaptation**,  
 654    where paths become redundant. In contrast, we show how RaBiT's coupled training mecha-  
 655    nism structurally enforces an **error-correcting hierarchy** and is superior to other heuristic  
 656    solutions.

657    **Proposition 1** (Inter-Path Adaptation in Standard QAT). *In a Standard QAT scheme  
 658    where two paths ( $\hat{\mathbf{W}}_1, \hat{\mathbf{W}}_2$ ) are updated from their respective latent weights ( $\mathbf{W}_1, \mathbf{W}_2$ ) using  
 659    a shared global gradient  $\mathbf{G} = \nabla_{\mathbf{W}_1 + \mathbf{W}_2} \mathcal{L}$ , the paths have a persistent tendency to become  
 660    positively correlated, leading to redundancy.*

661  
 662    *Proof.* Let the latent weights be  $\mathbf{W}_1$  and  $\mathbf{W}_2$ . After a single update step with learning rate  
 663     $\eta$  and shared gradient  $\mathbf{g}$ , the new weights are  $\mathbf{W}'_1$  and  $\mathbf{W}'_2$ :

$$\mathbf{W}'_1 := \mathbf{W}_1 - \eta \mathbf{G} \quad \text{and} \quad \mathbf{W}'_2 := \mathbf{W}_2 - \eta \mathbf{G}$$

664    The change in the Frobenius inner product between the weights, which reflects their corre-  
 665    lation, is:

$$\Delta_{\langle \cdot, \cdot \rangle} := \langle \mathbf{W}'_1, \mathbf{W}'_2 \rangle_F - \langle \mathbf{W}_1, \mathbf{W}_2 \rangle_F$$

666    Expanding this gives:

$$\Delta_{\langle \cdot, \cdot \rangle} = -\eta (\langle \mathbf{W}_1, \mathbf{G} \rangle_F + \langle \mathbf{W}_2, \mathbf{G} \rangle_F) + \eta^2 \|\mathbf{G}\|_F^2$$

667    While the linear terms depend on the alignment between the current weights and the gradient,  
 668    the quadratic term  $\eta^2 \|\mathbf{G}\|_F^2$  is **always non-negative**. This term acts as a systematic force,  
 669    constantly pushing the two paths in the same direction defined by the global gradient  $\mathbf{G}$ .  
 670    This dynamic, the underlying mechanism of **inter-path adaptation**, compels both paths  
 671    to learn redundant, dominant features in order to minimize the global loss. This ultimately  
 672    leads to a break down of the intended residual hierarchy and compromises the model's  
 673    expressive capacity.  $\square$

674    **Proposition 2** (Structurally Enforced Error Correction in RaBiT). *RaBiT's coupled training  
 675    mechanism resolves the redundancy drift by fundamentally changing the optimization objective.  
 676    Instead of independent updates, RaBiT's on-the-fly derivation structurally forces the second  
 677    path ( $\hat{\mathbf{W}}_2$ ) to align with the true residual of the first path ( $\mathbf{R}_1 = \mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_1$ ), thereby  
 678    enforcing an error-correcting relationship.*

679    *Analysis.* The optimization objectives of the two paths are implicitly different in RaBiT  
 680    versus the naïve approach.

- **Standard QAT Objective:** Both paths are driven by the same structurally-agnostic  
 681    global gradient  $\mathbf{G}$ . Their implicit goal is to align with  $\mathbf{G}$  to reduce the global loss. Since  
 682    both  $\hat{\mathbf{W}}_1$  and  $\hat{\mathbf{W}}_2$  are incentivized to align with the same vector  $\mathbf{G}$ , they inevitably learn  
 683    to align with each other, leading to redundancy as shown in Proposition 1.

$$\hat{\mathbf{W}}_1 \propto \mathbf{G} \quad \text{and} \quad \hat{\mathbf{W}}_2 \propto \mathbf{G} \implies \langle \hat{\mathbf{W}}_1, \hat{\mathbf{W}}_2 \rangle_F > 0$$

- **RaBiT's Enforced Objective:** RaBiT maintains a single shared blueprint,  $\mathbf{W}_{\text{FP}}$ . The  
 691    on-the-fly derivation process,  $\mathbf{R}_1 := \mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_1$  followed by the binarization of  $\mathbf{R}_1$  to  
 692    create  $\hat{\mathbf{W}}_2$ , explicitly defines the optimization target for the second path. The goal for  
 693     $\hat{\mathbf{W}}_2$  is no longer to align with the global gradient  $\mathbf{g}$ , but to be the best possible low-rank  
 694    approximation of the current residual  $\mathbf{R}_2$ .

$$\text{Objective for } \hat{\mathbf{W}}_2 : \min \|\mathbf{R}_1 - \hat{\mathbf{W}}_2\|_F^2 \implies \hat{\mathbf{W}}_2 \approx \mathbf{R}_1$$

695    This structural constraint forces a high **Residual Alignment**. In the context of extreme  
 696    low-bit quantization, the first approximation  $\hat{\mathbf{W}}_1$  often "overshoots" the target  $\mathbf{W}_{\text{FP}}$   
 697    in certain directions. To correct this, the residual  $\mathbf{R}_1 = \mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_1$  will point in the

702 opposite direction of the overshoot. By aligning with  $\mathbf{R}_1$ ,  $\hat{\mathbf{W}}_2$  naturally becomes **anti-**  
 703 **correlated** with  $\hat{\mathbf{W}}_1$ , implementing an efficient **active cancellation** mechanism rather  
 704 than degenerating into redundancy.  
 705  $\square$

706  
 707 **Proposition 3** (Superior Optimization Dynamics of Coupled vs. Iterative Training).  
 708 *Iterative training (e.g., freezing one path while training the other) avoids adaptation but at  
 709 the cost of optimization efficiency. In contrast, RaBiT resolves adaptation while permitting  
 710 full parameter co-adaptation, resulting in a superior optimization trajectory.*  
 711

712 *Proof.* Following Proposition 1, the problem of Standard QAT is the simultaneous update  
 713 of both paths in the same direction. An alternative solution is to update them iteratively,  
 714 which prevents this simultaneous push and thus avoids adaptation. However, this introduces  
 715 a new problem of inefficiency.  
 716

717 The optimal direction to reduce the loss  $\mathcal{L}$  is the steepest descent direction in the joint  
 718 parameter space of  $(\mathbf{W}_1, \mathbf{W}_2)$ , which is  $\mathbf{d}^* = (-\mathbf{G}, -\mathbf{G})$ . When training iteratively, one  
 719 path is frozen, so the update is restricted to an axis-aligned direction, e.g.,  $\mathbf{d}_{\text{iter}} = (\mathbf{0}, -\mathbf{G})$ .  
 720 The cosine similarity between the iterative update and the optimal update direction is:  
 721

$$\begin{aligned}\cos(\theta) &= \frac{\langle \mathbf{d}_{\text{iter}}, \mathbf{d}^* \rangle_F}{\|\mathbf{d}_{\text{iter}}\|_F \|\mathbf{d}^*\|_F} = \frac{\langle (\mathbf{0}, -\mathbf{G}), (-\mathbf{G}, -\mathbf{G}) \rangle_F}{\|(\mathbf{0}, -\mathbf{G})\|_F \|(-\mathbf{G}, -\mathbf{G})\|_F} \\ &= \frac{\|\mathbf{G}\|_F^2}{\|\mathbf{G}\|_F \cdot \sqrt{\|\mathbf{G}\|_F^2 + \|\mathbf{G}\|_F^2}} = \frac{1}{\sqrt{2}}\end{aligned}$$

722 This fixed  $45^\circ$  misalignment forces the optimization to follow an inefficient zig-zag trajectory.  
 723 While it solves adaptation, it sacrifices optimization efficiency.  
 724  $\square$

725 RaBiT, through its coupled derivation described in Proposition 2, resolves this trade-off.  
 726 By updating a single shared weight  $\mathbf{W}_{\text{FP}}$  with the full gradient  $\mathbf{G}$ , it allows both paths to  
 727 co-adapt simultaneously in a coordinated manner that is not restricted to an inefficient path.  
 728 Thus, RaBiT resolves adaptation without compromising optimization efficiency, leading to  
 729 superior dynamics.  $\square$

730 **Corollary 1 (to Proposition 2). Negative Correlation Induction.** *RaBiT’s coupled  
 731 training mechanism, by forcing the second path ( $\hat{\mathbf{W}}_2$ ) to approximate the residual of the  
 732 first path ( $\mathbf{R}_1$ ), inherently promotes a negative correlation between their respective outputs  
 733 ( $\mathbf{y}_1, \mathbf{y}_2$ ).*

734 *Analysis.* From Proposition 2, we established that RaBiT trains the second path to approximate  
 735 the residual of the first:  
 736

$$\hat{\mathbf{W}}_2 \approx \mathbf{R}_1 = \mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_1$$

737 Let’s consider the outputs for a given input  $\mathbf{x}$ . The outputs of the full-precision teacher, the  
 738 first path, and the second path are  $\mathbf{y}_t = \mathbf{W}_{\text{FP}}\mathbf{x}$ ,  $\mathbf{y}_1 = \hat{\mathbf{W}}_1\mathbf{x}$ , and  $\mathbf{y}_2 = \hat{\mathbf{W}}_2\mathbf{x}$ , respectively.  
 739 Based on the weight approximation, the output of the second path is:  
 740

$$\mathbf{y}_2 \approx \mathbf{R}_1\mathbf{x} = (\mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_1)\mathbf{x} = \mathbf{y}_t - \mathbf{y}_1$$

741 Now, we can analyze the covariance between the outputs  $\mathbf{y}_1$  and  $\mathbf{y}_2$ . Assuming the outputs  
 742 are centered for simplicity, the covariance is proportional to the expected value of their dot  
 743 product,  $\mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_2]$ .  
 744

$$\mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_2] \approx \mathbb{E}[\mathbf{y}_1^\top (\mathbf{y}_t - \mathbf{y}_1)] = \mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_t] - \mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_1] = \mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_t] - \mathbb{E}[\|\mathbf{y}_1\|^2]$$

745 Let’s analyze the two terms:  
 746

- 747 1.  $\mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_t]$ : The first path  $\hat{\mathbf{W}}_1$  is the primary, albeit coarse, approximation of  $\mathbf{W}_{\text{FP}}$ .  
 748 Its purpose is to capture the main features of the teacher. Therefore, their outputs  
 749  $\mathbf{y}_1$  and  $\mathbf{y}_t$  are expected to be strongly and **positively correlated**, making this term  
 750 a large positive value.

- 756      2.  $\mathbb{E}[\|\mathbf{y}_1\|^2]$ : This is the expected squared norm of the first path’s output. Binarization  
 757      is an aggressive quantization that often leads to an “overshoot” in magnitude. A  
 758      single binary path must represent a wide range of continuous values, so its effective  
 759      scaling factor often results in an output magnitude  $\|\mathbf{y}_1\|$  that exceeds the projection  
 760      of  $\mathbf{y}_1$  onto  $\mathbf{y}_t$ . Consequently, it is generally the case that  $\|\mathbf{y}_1\|^2 > \mathbf{y}_1^\top \mathbf{y}_t$ , making  
 761       $\mathbb{E}[\|\mathbf{y}_1\|^2]$  a larger positive term than  $\mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_t]$ .

762      Combining these points, the covariance is approximately the difference between a positive  
 763      term and a larger positive term:

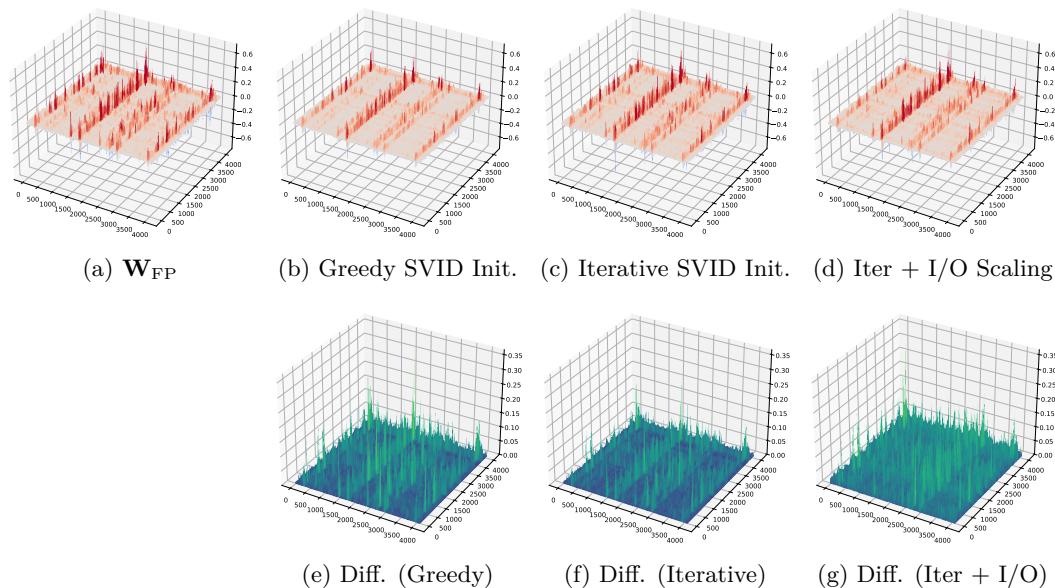
$$\text{Cov}(\mathbf{y}_1, \mathbf{y}_2) \approx \underbrace{\mathbb{E}[\mathbf{y}_1^\top \mathbf{y}_t]}_{\text{Positive Alignment}} - \underbrace{\mathbb{E}[\|\mathbf{y}_1\|^2]}_{\text{Larger Magnitude Term}} < 0$$

767      Thus, RaBiT’s mechanism of forcing the second path to correct the error of the first  
 768      path structurally drives the covariance, and therefore the correlation  $\text{Corr}(\mathbf{y}_1, \mathbf{y}_2)$ , to be  
 769      negative.  $\square$

## 771      A.2 INITIALIZATION ANALYSIS: FUNCTIONALITY VS. APPROXIMATION

773      Table 6: **Initialization Analysis on Llama2-7B.** Trade-off between weight reconstruction  
 774      error (Avg. MAE/MSE) and model functionality (Initial KL Divergence Loss), for the first  
 775      `q_proj` layer. I/O Channel Importance Scaling dramatically reduces KL Divergence Loss  
 776      despite increasing MSE.

Initialization Method	Avg. MAE ↓	Avg. MSE ↓	KL Loss ↓
Greedy SVID	0.359	0.150	17,152
Iterative Residual SVID + I/O Ch. Importance Scaling	0.370	0.122	13,760
	<b>0.632</b>	<b>0.302</b>	<b>2,672</b>



801      Figure 3: **Visual analysis of weight initialization for the first layer’s `q_proj` matrix**  
 802      **of the Llama2-7B model.** The top row displays the original full-precision weight ( $\mathbf{W}_{\text{FP}}$ )  
 803      alongside its initial approximations from three methods: (b) Greedy SVID,  
 804      (c) Iterative SVID, and (d) Iterative SVID with I/O Channel Importance Scaling.  
 805      The bottom row shows the corresponding difference matrices ( $\mathbf{W}_{\text{FP}} - \hat{\mathbf{W}}_{\text{init}}$ ), illustrating the initial error structure.  
 806      Our function-aware initialization produces a qualitatively different structure compared to  
 807      the others, which is reflected in its distinct error pattern.

808      Stable initialization is paramount in the low-bit regime, as the initial quantization er-  
 809      ror spike can destabilize QAT. On the Llama2-7b model, we evaluate our proposed tech-

niques—Iterative Residual SVID and I/O Channel Importance Scaling (Section 4.3)—by measuring both the weight reconstruction error (Avg. MAE, MSE) and the initial task loss (Knowledge Distillation (KD) loss) before the first training step.

Table 6 details the results averaged across attention projection layers and reveals a crucial insight. As our baseline, Greedy SVID is a non-iterative decomposition that finalizes each path sequentially without the co-adaptation enabled by our iterative approach. First, regarding **Iterative Refinement**, moving from Greedy SVID to Iterative Residual SVID consistently improves weight reconstruction (*e.g.*, Avg. MSE drops 0.150 → 0.122) and substantially reduces the initial KL divergence loss (17,152 → 13,760), confirming mitigation of scheduling bias. Second, adding **I/O Channel Importance Scaling** to the iterative process yields a striking result: while reconstruction error increases significantly (Avg. MSE 0.122 → 0.302), the KL divergence loss *plummets* dramatically (13,760 → 2,672, an 81% reduction).

This confirms that extreme quantization should prioritize preserving *functionality* over merely approximating *weights*. I/O Channel Importance Scaling allocates the limited 2-bit capacity to critical channels based on activation and gradient statistics (Section 4.3), sacrificing the reconstruction of less important weights. This trade-off is visually stark in Figure 3. While Iterative SVID produces a lower-error approximation than Greedy SVID (comparing Figure 3f to Figure 3e), the function-aware I/O Scaling method yields a visibly larger reconstruction error (Figure 3g). Despite this higher weight-level discrepancy, its focus on functional saliency provides a far superior starting point for QAT, as evidenced by the dramatic reduction in initial task loss.

### A.3 EXTENDED RESULTS

**Detailed Zero-Shot Reasoning Accuracy.** Table 7 and Table 8 provide a detailed breakdown of the zero-shot reasoning accuracy across five common benchmarks, complementing the average scores reported in the main text.

Table 7: **Detailed Zero-Shot Reasoning Accuracy on Llama Models (%)**. Comparison of FP16 against leading 2-bit methods on five common benchmarks.

Models	Method	WinoGrande↑	HellaSwag↑	ARC-e↑	ARC-c↑	PIQA↑	Average↑
Llama2-7B	FullPrecision	67.80	56.71	69.28	39.93	78.29	62.40
	QTIP	64.64	53.09	65.57	35.67	<b>75.90</b>	58.97
	DBF	63.61	52.44	64.73	35.58	75.84	58.44
	RaBiT (Ours)	<b>67.80</b>	<b>53.52</b>	<b>72.43</b>	<b>37.88</b>	<b>75.90</b>	<b>61.51</b>
Llama2-13B	FullPrecision	69.93	59.64	73.19	45.73	78.67	65.43
	QTIP	<b>67.56</b>	<b>57.4</b>	<b>70.8</b>	<b>41.46</b>	77.37	<b>62.92</b>
	DBF	67.09	56.6	69.02	38.74	78.18	61.93
	RaBiT (Ours)	<b>67.56</b>	56.71	69.06	39.76	<b>77.42</b>	62.10
Llama3-8B	FullPrecision	72.93	60.08	80.30	50.17	79.76	67.80
	QTIP	<b>70.24</b>	<b>55.53</b>	75.29	41.64	76.71	63.88
	DBF	68.90	54.49	74.62	39.76	76.44	62.84
	RaBiT (Ours)	69.37	55.13	<b>75.37</b>	<b>42.83</b>	<b>77.96</b>	<b>64.13</b>

Table 8: **Detailed Zero-Shot Reasoning Accuracy on Gemma Models (%)**. Comparison of FP16 against leading 2-bit methods on five common benchmarks.

Models	Method	WinoGrande↑	HellaSwag↑	ARC-e↑	ARC-c↑	PIQA↑	Average↑
Gemma3-1B	FullPrecision	59.59	47.30	72.22	35.32	74.65	57.82
	QTIP	54.62	38.24	63.93	25.85	68.88	50.30
	DBF	<b>58.01</b>	40.37	62.92	28.41	70.18	51.98
	RaBiT (Ours)	56.59	<b>42.94</b>	<b>64.52</b>	<b>29.44</b>	<b>72.42</b>	<b>53.18</b>
Gemma3-4B	FullPrecision	69.22	56.77	81.52	51.45	79.05	67.60
	QTIP	<b>66.85</b>	52.25	<b>77.53</b>	<b>44.62</b>	76.12	<b>63.47</b>
	DBF	63.69	50.15	74.74	40.87	75.08	60.91
	RaBiT (Ours)	65.19	<b>52.57</b>	75.04	41.38	<b>76.88</b>	62.21
Gemma3-12B	FullPrecision	75.45	61.98	87.08	61.60	81.12	73.45
	QTIP	<b>72.69</b>	57.99	<b>84.09</b>	<b>54.95</b>	78.73	<b>69.69</b>
	DBF	72.14	57.20	82.49	52.05	77.97	68.37
	RaBiT (Ours)	72.30	<b>58.45</b>	82.41	52.13	<b>78.95</b>	68.85

864     A.4 INFERENCE PERFORMANCE ANALYSIS  
 865  
 866       A.4.1 KERNEL DESIGN  
 867  
 868       Our CUDA kernels implement binary GEMV operations tailored to the memory-bound  
 869       regime typical of the decoding phase in LLM inference. The design centers on bit-packing to  
 870       reduce global memory traffic, with a latency-tolerant and matmul-free compute pipeline that  
 871       leverages register-level staging.  
 872  
 873       **Weight Packing.** To reduce memory traffic, each group of 32 columns is mapped to  
 874       a `uint32_t`, with  $+1 \mapsto 0$  and  $-1 \mapsto 1$ . We then group the 32-bit words into `uint2` or  
 875       `PackedBits3` ( $3 \text{ uint32\_t}$  weights with padding), for 2-bit (2 binary weights) and 3-bit (3  
 876       binary weights) models, respectively. Rows are interleaved into warp-sized groups, ensuring  
 877       that a warp issues full coalesced memory transactions when loading weights. Our efficient  
 878       packing reduces the raw footprint of weights by a factor of  $32\times$ , compared to full-precision  
 879       weights.  
 880  
 881       **Compute Pipeline.** Each warp is assigned a set of output rows to avoid inter-warp  
 882       synchronization. Input activations (`x`) and column scales (`g`) are read as vectorized `uint4`  
 883       chunks. Binary signs are applied via lane-local bit shifts and XOR masks, instead of matrix  
 884       multiplication. The kernel uses simple yet effective pipelining: while one tile of data is  
 885       consumed, the subsequent tile is prefetched into registers. Accumulation proceeds using  
 886       `half2` fused multiply-add intrinsics (`_hfma2`), which increase arithmetic throughput without  
 887       resorting to shared memory. Finally, reductions across threads in a warp are performed with  
 888       shuffle operations, and output scale factors (`h`) are applied in `fp16` precision. Notably, our  
 889       architecture enables per-path parallelizable computation - instead of an  $n$ -bit weight, we  
 890       parallelize with  $n$  1-bit operations.  
 891       Efficient weight packing and pipelining reduce global memory access and raise the utilization  
 892       of execution units on the GPU. The kernel therefore shifts the limiting factor from raw  
 893       memory bandwidth toward register throughput, yielding measurable efficiency gains during  
 894       the decoding stage of LLM inference, showing remarkable performance.  
 895  
 896       A.4.2 MORE COMPARISONS  
 897  
 898       To provide a more detailed analysis of the end-to-end inference speed, we benchmarked  
 899       RaBiT against several key baselines: the full-precision (FP16) model, QTIP as the state-of-  
 900       the-art Vector Quantization (VQ) method, and DBF, which features a similar stacked binary  
 901       architecture. For QTIP, we utilized the publicly available CUDA kernels from the official  
 902       implementation<sup>4</sup>. For DBF, which also uses a stacked binary design but executes its two  
 903       paths sequentially, we developed an optimized CUDA kernel that runs approximately 21%  
 904       faster than their public Triton-based implementation to ensure a fair and robust comparison.  
 905       All evaluations were conducted on an NVIDIA RTX 4090 using the Llama2-7B model.  
 906  
 907       The results, depicted in Figure 4, were benchmarked across a range of generated token lengths  
 908       (64, 128, 256, 512, and 1024) for a comprehensive analysis. As expected, all 2-bit methods  
 909       significantly outperform the FP16 baseline due to the  $8\times$  reduction in memory bandwidth  
 910       requirements. More importantly, **RaBiT demonstrates a substantial performance**  
 911       **advantage, achieving nearly twice the decoding throughput** of the other 2-bit  
 912       quantization methods. This speed-up stems directly from the efficiency of our parallel,  
 913       matmul-free architecture. Unlike DBF, which is bottlenecked by its sequential computation  
 914       of two binary paths, RaBiT’s fully parallel design allows it to maximally leverage the benefits  
 915       of its efficient binary cores. While the absolute tokens/second rate naturally decreases with  
 916       longer generation sequences, the relative performance gap between the methods remains  
 917       consistent, confirming the robustness of RaBiT’s architectural advantage.

<sup>4</sup><https://github.com/Cornell-RelaxML/qtip>

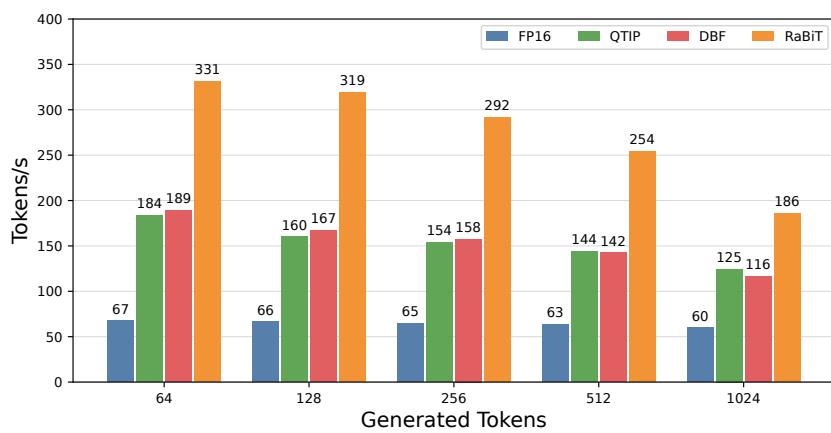


Figure 4: **End-to-end decoding throughput (tokens/second) for Llama2-7B on an NVIDIA RTX 4090 across various generated token lengths.** RaBiT’s parallel architecture consistently delivers superior performance over other 2-bit methods.

## A.5 HYPERPARAMETERS

### A.5.1 TRAINING DETAILS

We detail the hyperparameters used for our Quantization-Aware Training (QAT) experiments in Table 9. All models were trained for 6 epochs using the Muon optimizer (Jordan et al., 2024) with a cosine learning rate decay schedule. The models were initialized using our proposed function-aware strategy, with a fixed SVID iteration count of  $T_{\max} = 20$ . Key hyperparameters, such as the learning rate and the I/O Channel Importance Scaling intensities ( $\alpha_{\text{in}}, \alpha_{\text{out}}$ ), were fine-tuned for each specific model to achieve the best performance. All experiments were conducted on a single node equipped with four NVIDIA H100 GPUs.

Table 9: **RaBiT** Training Details

Bit	Target	Training Setup		Llama2		Llama3		Gemma3		
		7B	13B	8B	1B	4B	12B			
2	Intensities ( $\alpha_{\text{in}}, \alpha_{\text{out}}$ )	(0.8, 0.65)	(0.95, 0.45)	(0.85, 0.7)	(0.85, 0.7)	(0.95, 0.7)	(0.75, 0.6)			
	Iteration ( $T_{\max}$ )	20	20	20	20	20	20			
	Learning Rate	12e-6	1e-5	1e-5	1e-5	1e-5	5e-6			
	Epoch	6	6	6	6	6	6			
	# GPUs	1 × 4	1 × 4	1 × 4	1 × 4	1 × 4	1 × 4			
	# Training Hours	39	56	38	8	23	67			
3	Intensities ( $\alpha_{\text{in}}, \alpha_{\text{out}}$ )	(0.8, 0.65)	(0.95, 0.45)	(0.85, 0.7)	-	-	-			
	Iteration ( $T_{\max}$ )	20	20	20	-	-	-			
	Learning Rate	1e-5	1e-5	1e-5	-	-	-			
	Epoch	6	6	6	-	-	-			
	# GPUs	1 × 4	1 × 4	1 × 4	-	-	-			
	# Training Hours	46	88	44	-	-	-			

### A.5.2 GRID SEARCH FOR I/O CHANNEL IMPORTANCE SCALING INTENSITIES

To determine the optimal intensity hyperparameters for our I/O Channel Importance Scaling (Section 4.3), we performed a comprehensive grid search. The objective was to identify the values of  $\alpha_{\text{in}}$  and  $\alpha_{\text{out}}$  that minimized the initial Knowledge Distillation (KD) loss post-initialization. This process utilized a calibration dataset of 128 samples randomly selected from the training data to measure the loss.

The example results of this search on the Llama2-7B model are detailed in Table 10. We observed a clear optimum, with the minimum initial KL divergence loss of 2,672 achieved at the configuration of  $\alpha_{\text{in}} = 0.80$  and  $\alpha_{\text{out}} = 0.65$ . This finding underscores the importance of a

balanced preconditioning strategy that considers both input activation statistics and output gradient magnitudes. We repeated this grid search process for all other models to find their optimal alpha values.

Table 10: **Grid search results for I/O Channel Importance Scaling Intensities ( $\alpha_{in}$ ,  $\alpha_{out}$ ) on Llama2-7B.** The metric is the Initial KL Divergence Loss (Lower is better). The optimal configuration is highlighted in bold.

$\alpha_{out}$	$\alpha_{in}$			
	0.75	<b>0.80</b>	0.85	0.90
0.55	3,100	2,932	2,984	3,108
0.60	2,932	2,938	2,971	3,143
<b>0.65</b>	3,167	<b>2,672</b>	2,697	3,063
0.70	3,083	2,821	2,983	3,462

### A.5.3 SVID ITERATION CONVERGENCE ANALYSIS

The Iterative Residual SVID initialization (Section 4.3) aims to mitigate the scheduling bias inherent in standard greedy initialization. We analyzed the required number of iterations ( $T_{max}$ ) for convergence on the Llama2-7B model. We measured the Initial KL divergence loss as the iterations progressed from 1 (equivalent to Greedy SVID) up to 35.

The results, shown in Figure 5, indicate that the initialization quality improves rapidly in the initial phase. The loss stabilizes significantly around 15 iterations, and the optimum is reached at 20 iterations. Beyond this point, further iterations do not provide additional benefits. Based on this analysis, we selected  $T_{max} = 20$  as the default setting for RaBiT initialization, providing an optimal balance between initialization quality and computational cost.

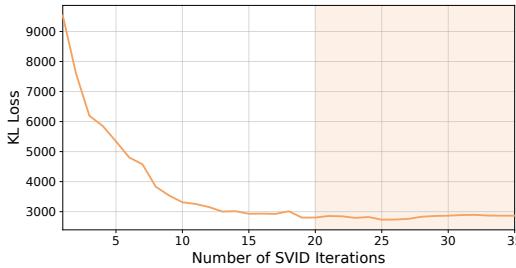


Figure 5: **Convergence analysis of Iterative Residual SVID on Llama2-7B.** The metric is the Initial KL Divergence Loss (Lower is better). Convergence stabilizes around 20 iterations.

### A.6 EXTENDED ANALYSIS OF INTER-PATH ADAPTATION

To provide a more granular view of the training dynamics, we conduct a layer-wise analysis of the Mean Squared Error (MSE) decomposition for the Llama2-7B model, visualized in Figure 6. This analysis offers two key insights into RaBiT’s structural advantages over Standard QAT.

First, the results empirically confirm our central hypothesis across the network’s depth. For most layers, RaBiT consistently generates a substantial negative covariance (the red-dashed component), which acts as a significant loss-reducing bonus, thereby lowering the total MSE. In contrast, Standard QAT fails to establish this effective error-cancellation, exhibiting a much smaller covariance term that provides negligible benefit. This provides strong visual evidence that RaBiT’s coupled training successfully enforces the intended error-correction hierarchy, while Standard QAT suffers from the performance degradation of inter-path adaptation.

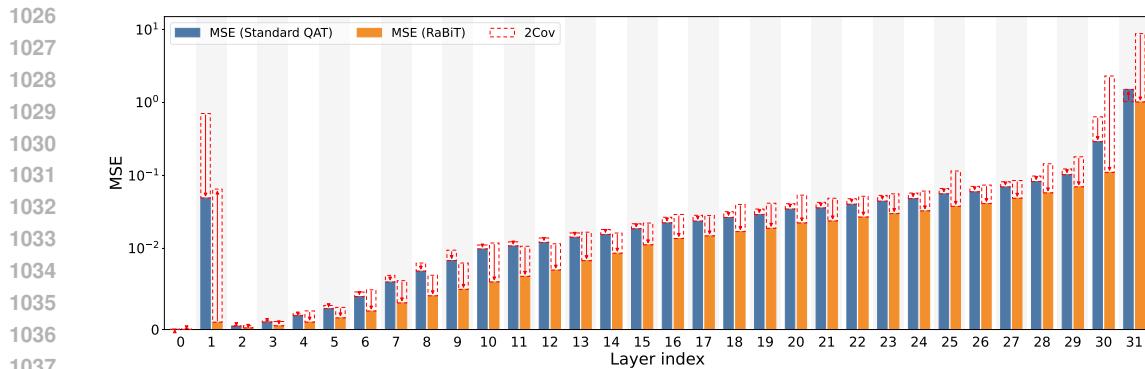


Figure 6: **Layer-wise MSE Decomposition in Llama2-7B’s down\_proj layers.** The bars compare the total Mean Squared Error (MSE) for Standard QAT (blue) and RaBiT (orange). The top of each bar represents the base error term ( $C$ ), while the red-dashed component visualizes twice the covariance ( $2 \times \text{Cov}$ ). RaBiT consistently generates a large negative covariance, which actively reduces the total MSE, demonstrating effective error cancellation. Notably, RaBiT also suppresses the extremely high MSE peak observed in the early layers of the Standard QAT baseline, indicating its robustness against layer sensitivity.

Second, and more strikingly, the analysis reveals RaBiT’s ability to overcome a critical optimization challenge in extreme quantization: layer sensitivity. The Standard QAT baseline exhibits an exceptionally high MSE in the initial layers, a phenomenon consistent with the known sensitivity of early network layers to input distributions and quantization errors, as also observed in other LLMs by (Zhang et al., 2025; 2024). RaBiT, however, dramatically suppresses this MSE peak. This suggests its benefits extend beyond merely enforcing anti-correlation. The fact that RaBiT tames this instability indicates that our method may resolve a more fundamental bottleneck in MSE-based QAT that has historically hindered extreme quantization in conventional architectures. While we designed RaBiT to foster negative correlation, its success in stabilizing these sensitive layers points to a deeper robustness. A full investigation into how residual coupling imparts this stability is a compelling direction for future research.

1080 A.7 GENERATED SAMPLES  
 1081

1082 SAMPLE 1: NEWTON  
 1083

1084 **Prompt:** Newton discovered the law of universal gravitation. Universal gravitation is  
 1085

---

1086 Generated Output  
 1087

1088 Newton discovered the law of universal gravitation. Universal gravitation is the force of  
 1089 gravity which draws objects towards the earth and other massive bodies. Gravity is the  
 1090 universal attraction between two objects and is proportional to the product of the masses  
 1091 of the two objects and inversely proportional to the square of the distance between them.  
 1092 Newton explained the force of gravity as being caused by the presence of matter in the  
 1093 universe. He stated that all matter attracted all other matter and that the tendency to  
 1094 move towards a massive body was proportional to the product of the masses of the objects.  
 1095 Newton developed the law of universal gravitation so that the attraction between two bodies  
 1096 could be calculated.

---

1096

1097

1098

1099 SAMPLE 2: ADAM SMITH  
 1100

1101 **Prompt:** Adam Smith wrote The Wealth of Nations. The Wealth of Nations is  
 1102

---

1103 Generated Output  
 1104

1105 Adam Smith wrote The Wealth of Nations. The Wealth of Nations is the first book to outline  
 1106 the modern principles of capitalism and to make the case for it. Smith's work is widely  
 1107 considered to be the first comprehensive system of political economy and one of the most  
 1108 important works of political philosophy. The Wealth of Nations is the foundation of modern  
 1109 capitalism. The Wealth of Nations is a book by Adam Smith, first published in 1776, which  
 1110 describes the nature of economic systems. Adam Smith is regarded as the father of modern  
 1111 economics. The Wealth of Nations is widely considered to be the first comprehensive system  
 1112 of political economy and one of the most important works of political philosophy. It is also  
 1113 one of the most influential books on economics. Adam Smith is regarded as the father of  
 1114 modern economics.

---

1114

1115

1116 SAMPLE 3: DNA DOUBLE HELIX  
 1117

1118 **Prompt:** Watson and Crick discovered the DNA double helix. The DNA double helix is  
 1119

---

1120 Generated Output  
 1121

1122 Watson and Crick discovered the DNA double helix. The DNA double helix is a structure  
 1123 found in the nucleus of every living cell. The double helix is the most fundamental structure  
 1124 in DNA. The structure is made of two single strands of DNA. The double helix was discovered  
 1125 in 1953 by James Watson and Francis Crick. In 1953, James Watson and Francis Crick  
 1126 discovered the double helix structure of DNA. The double helix structure is the most  
 1127 fundamental structure in DNA. The double helix structure is made of two single strands of  
 1128 DNA. The double helix is found in the nucleus of every living cell.

---

1128

1129

1130

1131

1132

1133

1134    A.8 ALGORITHMS  
 1135

---

**Algorithm 1** RaBiT Initialization: Calibrated Iterative Residual SVID
 

---

```

1: Require: Pretrained weight  $\mathbf{W}_{\text{FP}}$ , Number of paths  $k$ , Max iterations  $T_{\max}$ 
2: Require: Calibration stats  $\mathbf{s}_{\text{in}}, \mathbf{s}_{\text{out}}$  and intensities  $\alpha_{\text{in}}, \alpha_{\text{out}}$ 
3: Output: Initialized scales  $\{(\mathbf{g}_i, \mathbf{h}_i)\}_{i=1}^k$ 

4: // Step 1: I/O Channel Importance-Calibrated Preconditioning
5: Normalize:  $\mathbf{s}_{\text{in}} \leftarrow \mathbf{s}_{\text{in}} / \max(\mathbf{s}_{\text{in}})$ ,  $\mathbf{s}_{\text{out}} \leftarrow \mathbf{s}_{\text{out}} / \max(\mathbf{s}_{\text{out}})$ 
6: Precondition:  $\mathbf{W}' \leftarrow \mathbf{s}_{\text{out}}^{\alpha_{\text{out}}} \odot \mathbf{W}_{\text{FP}} \odot \mathbf{s}_{\text{in}}^{\alpha_{\text{in}}}$ 

7: // Step 2: Iterative Residual SVID
8: Initialize  $\hat{\mathbf{W}}_i^{(0)} \leftarrow \mathbf{0}$  for  $i = 1, \dots, k$ 
9: for  $t = 1$  to  $T_{\max}$  do
10:    for  $i = 1$  to  $k$  do
11:     // Calculate target residual (Gauss-Seidel style update)
12:      $\mathbf{R}_i^{(t)} \leftarrow \mathbf{W}' - \left( \sum_{j < i} \hat{\mathbf{W}}_j^{(t)} + \sum_{j > i} \hat{\mathbf{W}}_j^{(t-1)} \right)$ 
13:     // Apply SVID to find the best rank-1 approximation
14:      $(\mathbf{B}_i^{(t)}, \mathbf{g}_i^{(t)}, \mathbf{h}_i^{(t)}) \leftarrow \text{SVID}(\mathbf{R}_i^{(t)})$ 
15:      $\hat{\mathbf{W}}_i^{(t)} \leftarrow \mathbf{g}_i^{(t)} \odot \mathbf{B}_i^{(t)} \odot \mathbf{h}_i^{(t)}$ 
16:    end for
17: end for

18: // Step 3: Map scales back to the original weight domain
19: for  $i = 1$  to  $k$  do
20:     $\mathbf{g}_i \leftarrow \mathbf{s}_{\text{out}}^{-\alpha_{\text{out}}} \odot \mathbf{g}_i^{(T_{\max})}$ 
21:     $\mathbf{h}_i \leftarrow \mathbf{s}_{\text{in}}^{-\alpha_{\text{in}}} \odot \mathbf{h}_i^{(T_{\max})}$ 
22: end for

```

---

1159

---

**Algorithm 2** RaBiT: Residual-Aware Binarization Training (One Step)
 

---

```

1: Parameters: Shared full-precision weight  $\mathbf{W}_{\text{FP}}$ ; Scales  $\{(\mathbf{g}_i, \mathbf{h}_i)\}_{i=1}^k$ .
2: Input: Minibatch Input  $\mathbf{X}$ , Targets  $\mathbf{T}$ .
3: // 1. Forward Pass: On-the-fly Residual Coupling
4:  $\mathbf{R}_0 \leftarrow \mathbf{W}_{\text{FP}}$ . // Initialize residual with the shared weight
5:  $\hat{\mathbf{W}}^{(k)} \leftarrow \mathbf{0}$ . // Effective weight for the entire layer
6: for  $i = 1$  to  $k$  do
7:    // Sequentially derive the  $i$ -th binary path
8:     $\mathbf{B}_i \leftarrow \text{sign}(\mathbf{R}_{i-1})$ .
9:     $\hat{\mathbf{W}}_i \leftarrow \mathbf{g}_i \odot \mathbf{B}_i \odot \mathbf{h}_i$ .
10:     $\hat{\mathbf{W}}^{(k)} \leftarrow \hat{\mathbf{W}}^{(k)} + \hat{\mathbf{W}}_i$ .
11:    // Update residual for the next path
12:     $\mathbf{R}_i \leftarrow \mathbf{R}_{i-1} - \hat{\mathbf{W}}_i$ .
13: end for
14:  $\mathbf{Y} \leftarrow \hat{\mathbf{W}}^{(k)} \mathbf{X}$ . // Compute layer output
15: Calculate Loss  $\mathcal{L}(\mathbf{Y}, \mathbf{T})$ .

16: // 2. Backward Pass
17:  $\Delta \leftarrow \partial \mathcal{L} / \partial \mathbf{Y}$ . (Output gradient)
18: // Surrogate gradient for the shared weight  $\mathbf{W}_{\text{FP}}$ 
19:  $\nabla_{\mathbf{W}_{\text{FP}}} \leftarrow \Delta^\top \mathbf{X}$ .
20: // Gradients for scales (treating  $\mathbf{B}_i$  as constant)
21: for  $i = 1$  to  $k$  do
22:    Compute  $\nabla_{\mathbf{g}_i}$  and  $\nabla_{\mathbf{h}_i}$  using  $\Delta, \mathbf{B}_i, \mathbf{X}$ , and other scales.
23: end for
24: // 3. Parameter Update
25: Update  $\{\mathbf{W}_{\text{FP}}, (\mathbf{g}_i, \mathbf{h}_i)_{i=1}^k\}$  using an optimizer with the computed gradients.

```

---

1184  
 1185  
 1186  
 1187

1188 A.9 LLM USAGE  
1189

1190 LLM is used only for writing, editing, or formatting purposes and does not impact the core  
1191 methodology, scientific rigorousness, or originality of the research.  
1192  
1193  
1194  
1195  
1196  
1197  
1198  
1199  
1200  
1201  
1202  
1203  
1204  
1205  
1206  
1207  
1208  
1209  
1210  
1211  
1212  
1213  
1214  
1215  
1216  
1217  
1218  
1219  
1220  
1221  
1222  
1223  
1224  
1225  
1226  
1227  
1228  
1229  
1230  
1231  
1232  
1233  
1234  
1235  
1236  
1237  
1238  
1239  
1240  
1241