

Knowledge Distillation in Time-series

October 8, 2020
Seonyoung Kim

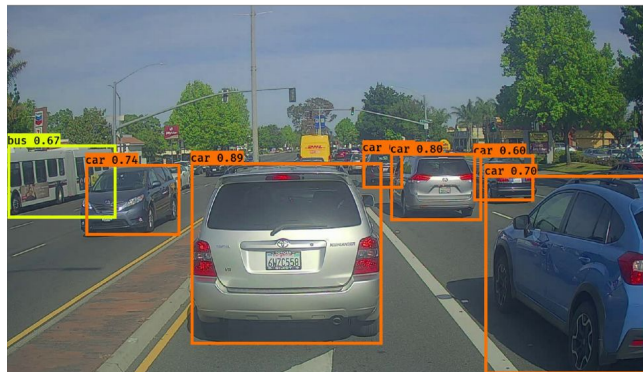
Model compression

Cloud computing



In the past, Cloud servers were used to run AI.

However, **speed, privacy, and cost problems** occurred, which is very important to some Applications.



Self driving cars



Security robots

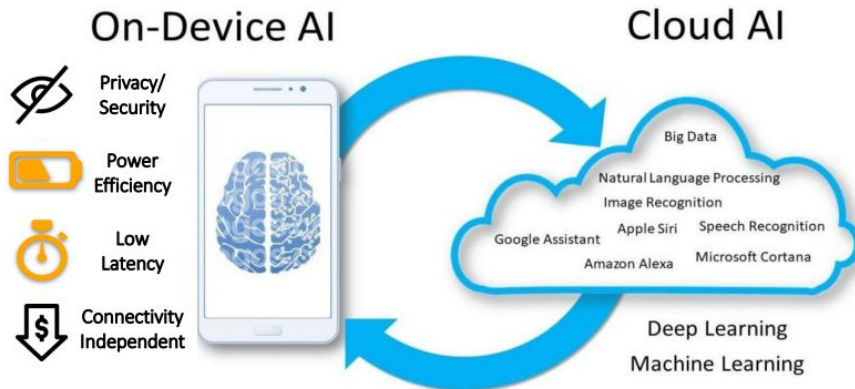
On-device AI



AI without cloud server.

However, it is impossible to use only on-device AI.

→ Model training is executed in cloud server & Pre-trained model is inference to device




Model compression


Need to **reduce the size of the model** for fast speed, small storage space, and low battery consumption.

→ Model compression technique

- Pruning, Quantization, Knowledge distillation, Low-rank approximation, Compact networks design




Original model on Cloud server




Compressed model on mobile device

Knowledge distillation(KD)



One of the techniques of model compression

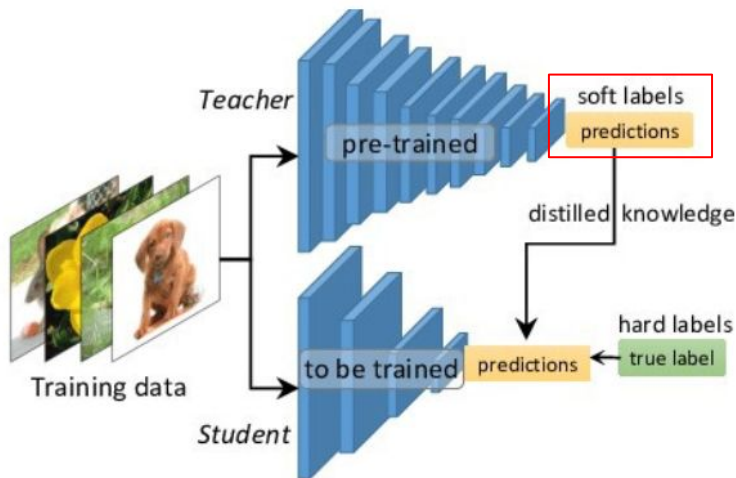
A method of distilling important knowledge of a **large neural network** and delivering it to a **small neural network**

distillation 미국식 [dɪstəˈleɪʃən]  [다른 뜻\(1건\)](#)

[U] 증류(법), [UC] 증류물, 정수

Knowledge distillation (KD)

Distilling the Knowledge in a Neural Network(Hinton Geoffrey, NIPS 2014 Workshop)



Teacher network's output(i.e., soft label) = Knowledge

cow	dog	cat	car
0	1	0	0
cow	dog	cat	car
10^{-6}	.9	.1	10^{-9}

Hard label

Soft label

By passing the output of the teacher network to the student network, the student network is trained on it.

Large Network → Teacher Network

Small Network → Student Network


Knowledge distillation in Time-series data

Knowledge distillation in Time-series data

Long-Term Prediction of Small Time-Series Data Using Generalized Distillation(Hayashi et al., IJCNN 2019)


Time-series prediction

1. Multi-step prediction



2. Long-term prediction

- Task to predict for the distant future
- e.g., whether the stock price rise after 3 months?



Knowledge distillation in Time-series

Long-Term Prediction of Small Time-Series Data Using Generalized Distillation(Hayashi et al., IJCNN 2019)

Motivation


- There exist area where **only limited amount of data are available**(e.g., medical experiments, experiment with a small budget and etc.) → Cold-start problem
- Need to train model from small data

⇒ It use Knowledge distillation for **efficient learning, not model compression**

⇒ No difference in the size of Teacher network and Student network

Knowledge distillation in Time-series

Long-Term Prediction of Small Time-Series Data Using Generalized Distillation(Hayashi et al., IJCNN 2019)





1. Train a teacher model using **privileged information**
2. Generate soft targets using the teacher model
3. Train a student model with a set of hard targets and a set of soft targets

Knowledge distillation in Time-series

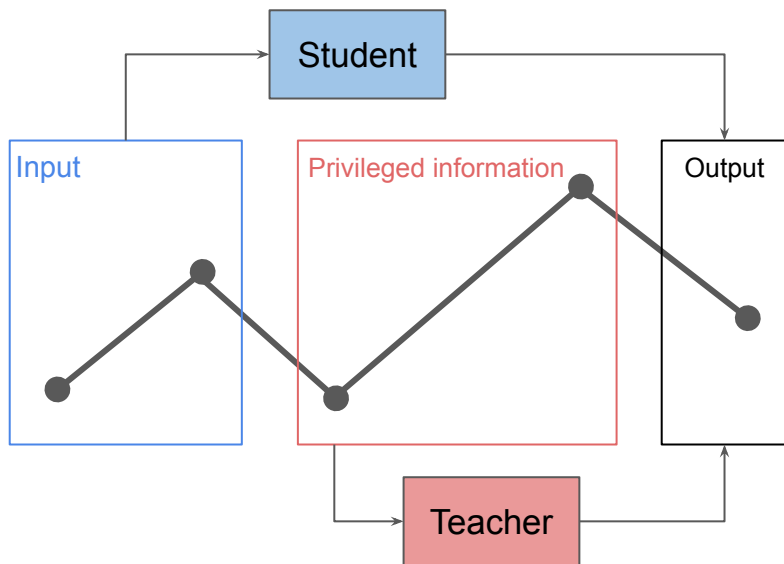
Privileged information

- *A new learning paradigm: Learning using privileged information*(V. Vapnik et al., 2009)
- Information is available at training time
- The selection of input & privileged information is arbitrary
- Example



Knowledge distillation in Time-series

Method



1. Train Teacher network using Privileged information
2. Train a student model with a set of hard targets and a set of soft targets

$$f_s = \arg \min_{f \in F_s} \frac{1}{n} \sum_{i=1}^n [(1 - \lambda)l(y_i, \sigma(f(x_i))) + \lambda l(s_i, \sigma(f(x_i)))]$$

f_s : student model

λ : imitation parameter $\in [0, 1]$

l : cross entropy loss function


s_i : soft target of Teacher model

Knowledge distillation in Time-series

Experiment

- Dataset : Mackey-Glass data, Beijing PM2.5 Data
- Logistic regression model
- Binary classification

Privileged Information



1. AI(data after the input)
2. M(data at mid-time between the input and the output)
3. BAO(data before and after the output)

Knowledge distillation in Time-series

Experiment (1) - Mackey-Glass Data

- synthetic data :
$$\frac{dx(t)}{dt} = -ax(t) + \frac{bx(t - t_0)}{1 + x(t - t_0)^c}$$
$$a = 0.0, b = 0.18, c = 10, t_0 = 16, x(t_0) = 0.9$$
- task : whether the k-step-ahead value(x_{t+k}) larger or smaller than the current one(x_t)?
- y value : binary-label based on k-step-ahead value

$$y_t = I(x_t \leq x_{t+k})$$

I : the indicator function

$$k = 8$$

- Example $x_0 = 0$ and $x_8 = 10$
then $y_0 = 1$

Knowledge distillation in Time-series

Experiment (2) - Beijing PM2.5 Data


- real data & multivariate time-series data
- Eight-dimensional data containing the concentration of dust in air
- Use only one feature
- task : whether the k-step-ahead PM2.5 value larger or smaller than the current one?
- y value : binary-label based on k-step-ahead value

$$y_t = I(x_t \leq x_{t+k})$$

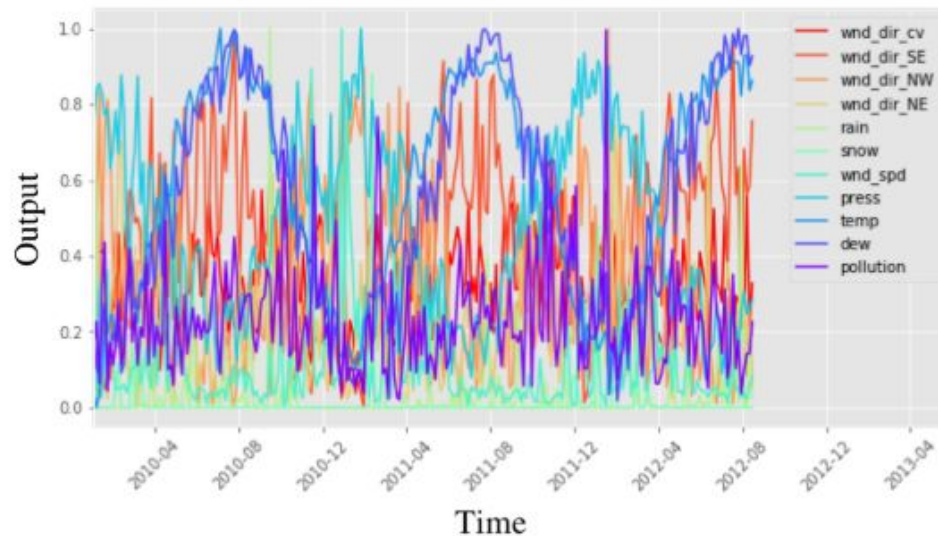
I : the indicator function

$$k = 15$$

Knowledge distillation in Time-series



Mackey-Glass Data



Beijing PM2.5 Data

Knowledge distillation in Time-series

Experiment Result

Data	Input length (student)	Input dim (student)	Input length (teacher)	Input dim (teacher)	Prediction time k	Training Time-series length	Validation Time-series length	Test time-series length
Mackey-Glass Data	4	4	2	2	8	100	100	10000
PM2.5-a Data	3	33	2	22	15	110	100	184
PM2.5-b Data	7	77	2	22	15	110	100	184

Table 1. Experimental setup

Method	Baseline (Logistic regression)	Proposed method
Mackey-Glass Data	0.971 (w/o regularizer)	0.932 (λ : 0.25, PI: BAO, T : 1.0, w/0 regularizer)
PM2.5-a Data	0.679 ($L1$: 1.0)	0.654 (λ : 0.5, PI: AI, T : 10.0, $L1$: 1.0)
PM2.5-b Data	0.690 ($L1$: 1.0)	0.716 (λ : 0.25, PI: BAO, T : 10.0, $L1$: 1.0)

Table 2. Experimental result

1. Mackey-Glass data
 - the problem might be too simple
2. PM2.5-a Data
 - the student input is too small

Knowledge distillation in Time-series

Conclusion

1. This paper proposed method about long-term prediction using privileged information
2. It showed that we should carefully select privileged information
3. It works well when data are multi-dimensional and hypothesis space is large
4. It need further experiment of more complex and difficult cases with higher-dimensional time-series data